

Basics on Conformance Checking

Josep Carmona

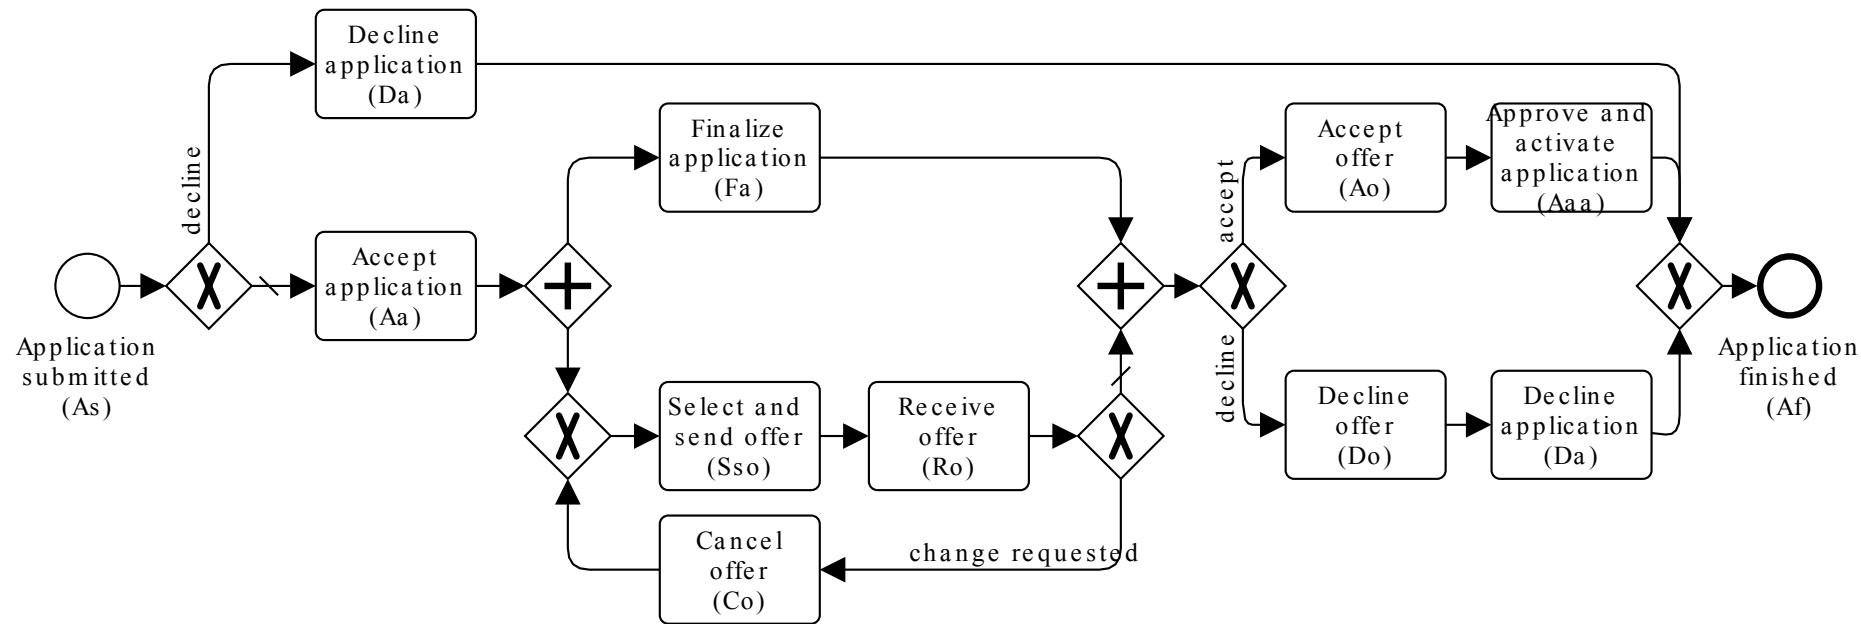
The Loan Application Process



A process defines how a loan application is handled once it has been submitted, with the outcome being its positive or negative assessment. A first activity is a check for eligibility of the respective applicant. Subsequently, a decision is taken, which determines whether an application is declined or accepted. If accepted, the application is finalised and an order is selected for further processing.

...

Process Models: BPMN



Control-flow

Data

Organizational

Other: Flexibility, ...

Covered Partially Covered Not Covered

Process Models (in Organizations)

- Process Models encompass knowledge in Organizations
- Unfortunately, they are not always:
 1. Known
 2. Understood
 3. Executed correctly
 4. Flexible enough when necessary, ...
- Organizations need to have a **systematic** way to be aware of these problems so that remedies can be put in place: **Conformance Checking**

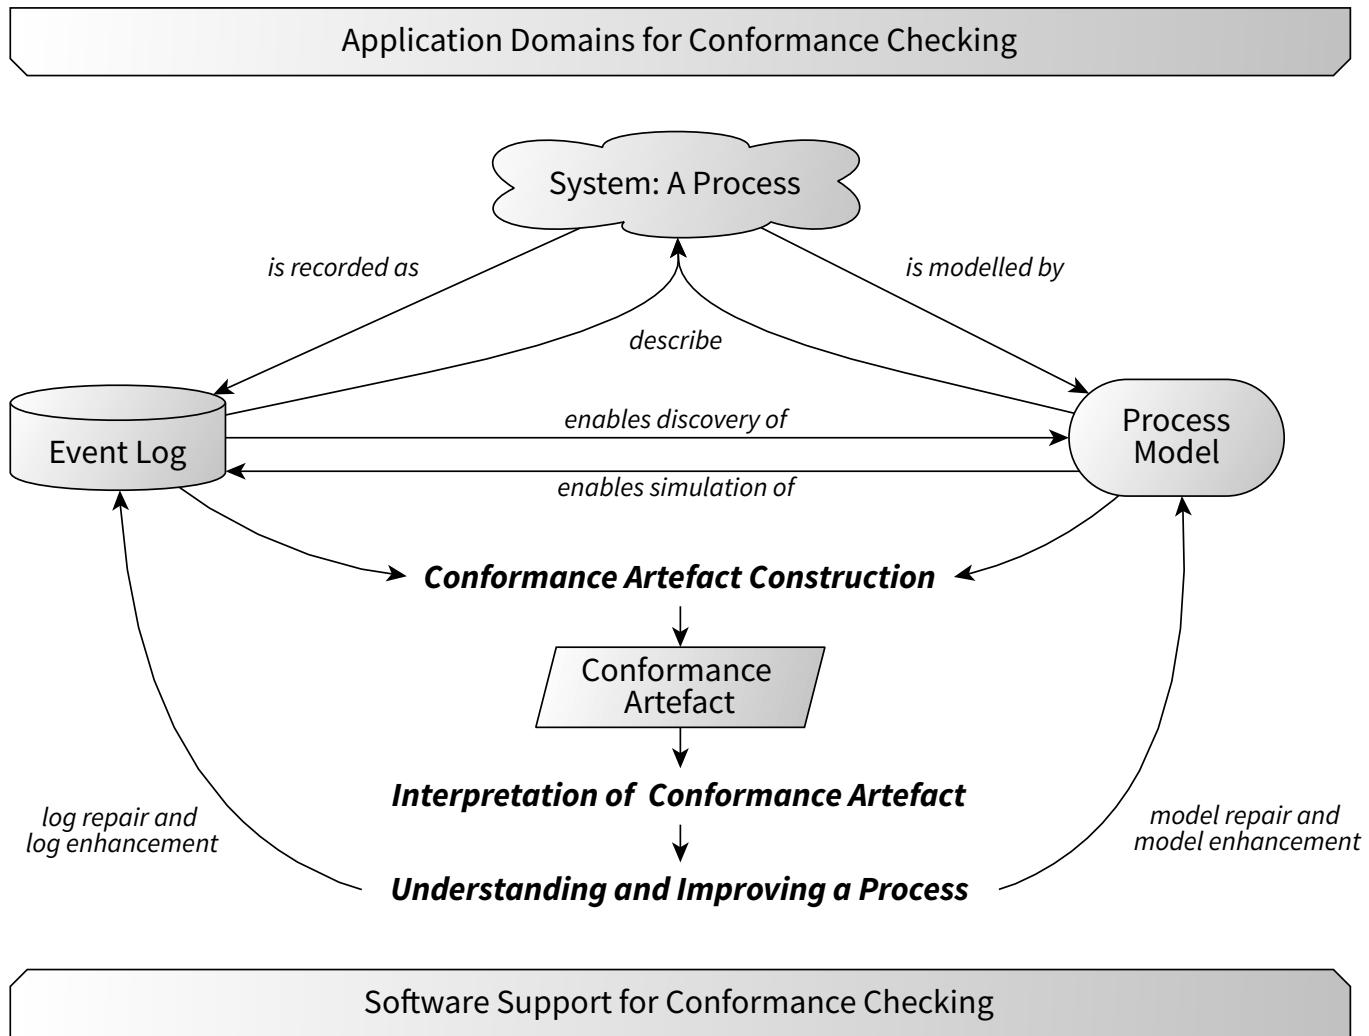
Event Logs

- Event log: recorded executions of a particular process (the reality!!!)
- Trace: Ordered collection of events corresponding to the same execution

Case A5634: < As, Aa, Fa, Sso, Ro, D, Da, Af >

Event	Application	Offer	Activity	Amount	Signed	Timestamp
...
<i>e₁₃</i>	A5634		Application submitted	€2,000		Jan 01, 12:31
<i>e₁₄</i>	A5634		Accept application	€2,000		Jan 01, 12:32
<i>e₁₅</i>	A5635		Application submitted	€5,000		Jan 02, 04:31
<i>e₁₆</i>	A5635		Accept application	€5,000		Jan 02, 04:32
<i>e₁₇</i>	A5636		Application submitted	€200		Jan 03, 06:59
<i>e₁₈</i>	A5636		Accept application	€200		Jan 03, 07:00
...
<i>e₂₂</i>	A5634		Finalise application			Jan 03, 09:00
<i>e₂₃</i>	A5636		Finalise application			Jan 03, 09:01
<i>e₂₄</i>	A5635		Decline application			Jan 03, 09:02
<i>e₂₅</i>	A5635		Decline application			Jan 03, 09:03
...
<i>e₃₀</i>	A5636	O3521	Select and send offer	€500		Jan 04, 16:32
...
<i>e₃₇</i>	A5634	O3541	Select and send offer	€1,500		Jan 05, 12:32
<i>e₃₈</i>	A5636	O3521	Receive offer		NO	Jan 05, 12:33
<i>e₃₈</i>	A5636	O3521	Cancel offer			Jan 05, 12:34
<i>e₃₉</i>	A5636	O3542	Select and send offer	€500		Jan 05, 13:29
<i>e₄₀</i>	A5636	O3542	Receive offer		YES	Jan 08, 08:33
<i>e₄₁</i>	A5636	O3542	Accept offer			Jan 08, 16:34
<i>e₄₂</i>	A5634	O3541	Receive offer		NO	Jan 10, 10:00
...
<i>e₅₄</i>	A5634	O3541	Decline offer			Jan 10, 10:04
...
<i>e₆₄</i>	A5634		Decline application			Jan 10, 10:05
<i>e₆₅</i>	A5634		Application finished			Jan 10, 10:06
<i>e₆₆</i>	A5636		Approve and activate application			Jan 10, 10:07
<i>e₆₇</i>	A5636		Application finished			Jan 10, 10:08
...

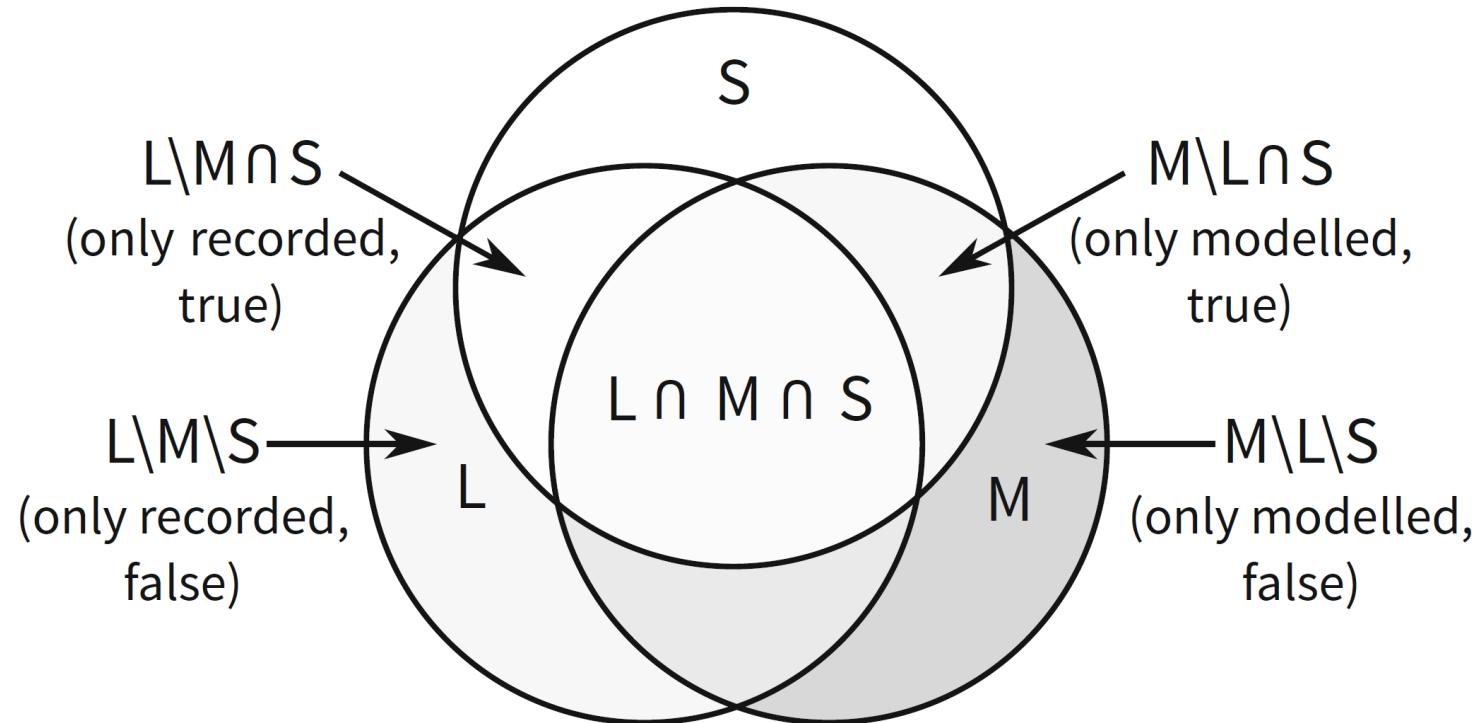
The Spectrum of Conformance Checking



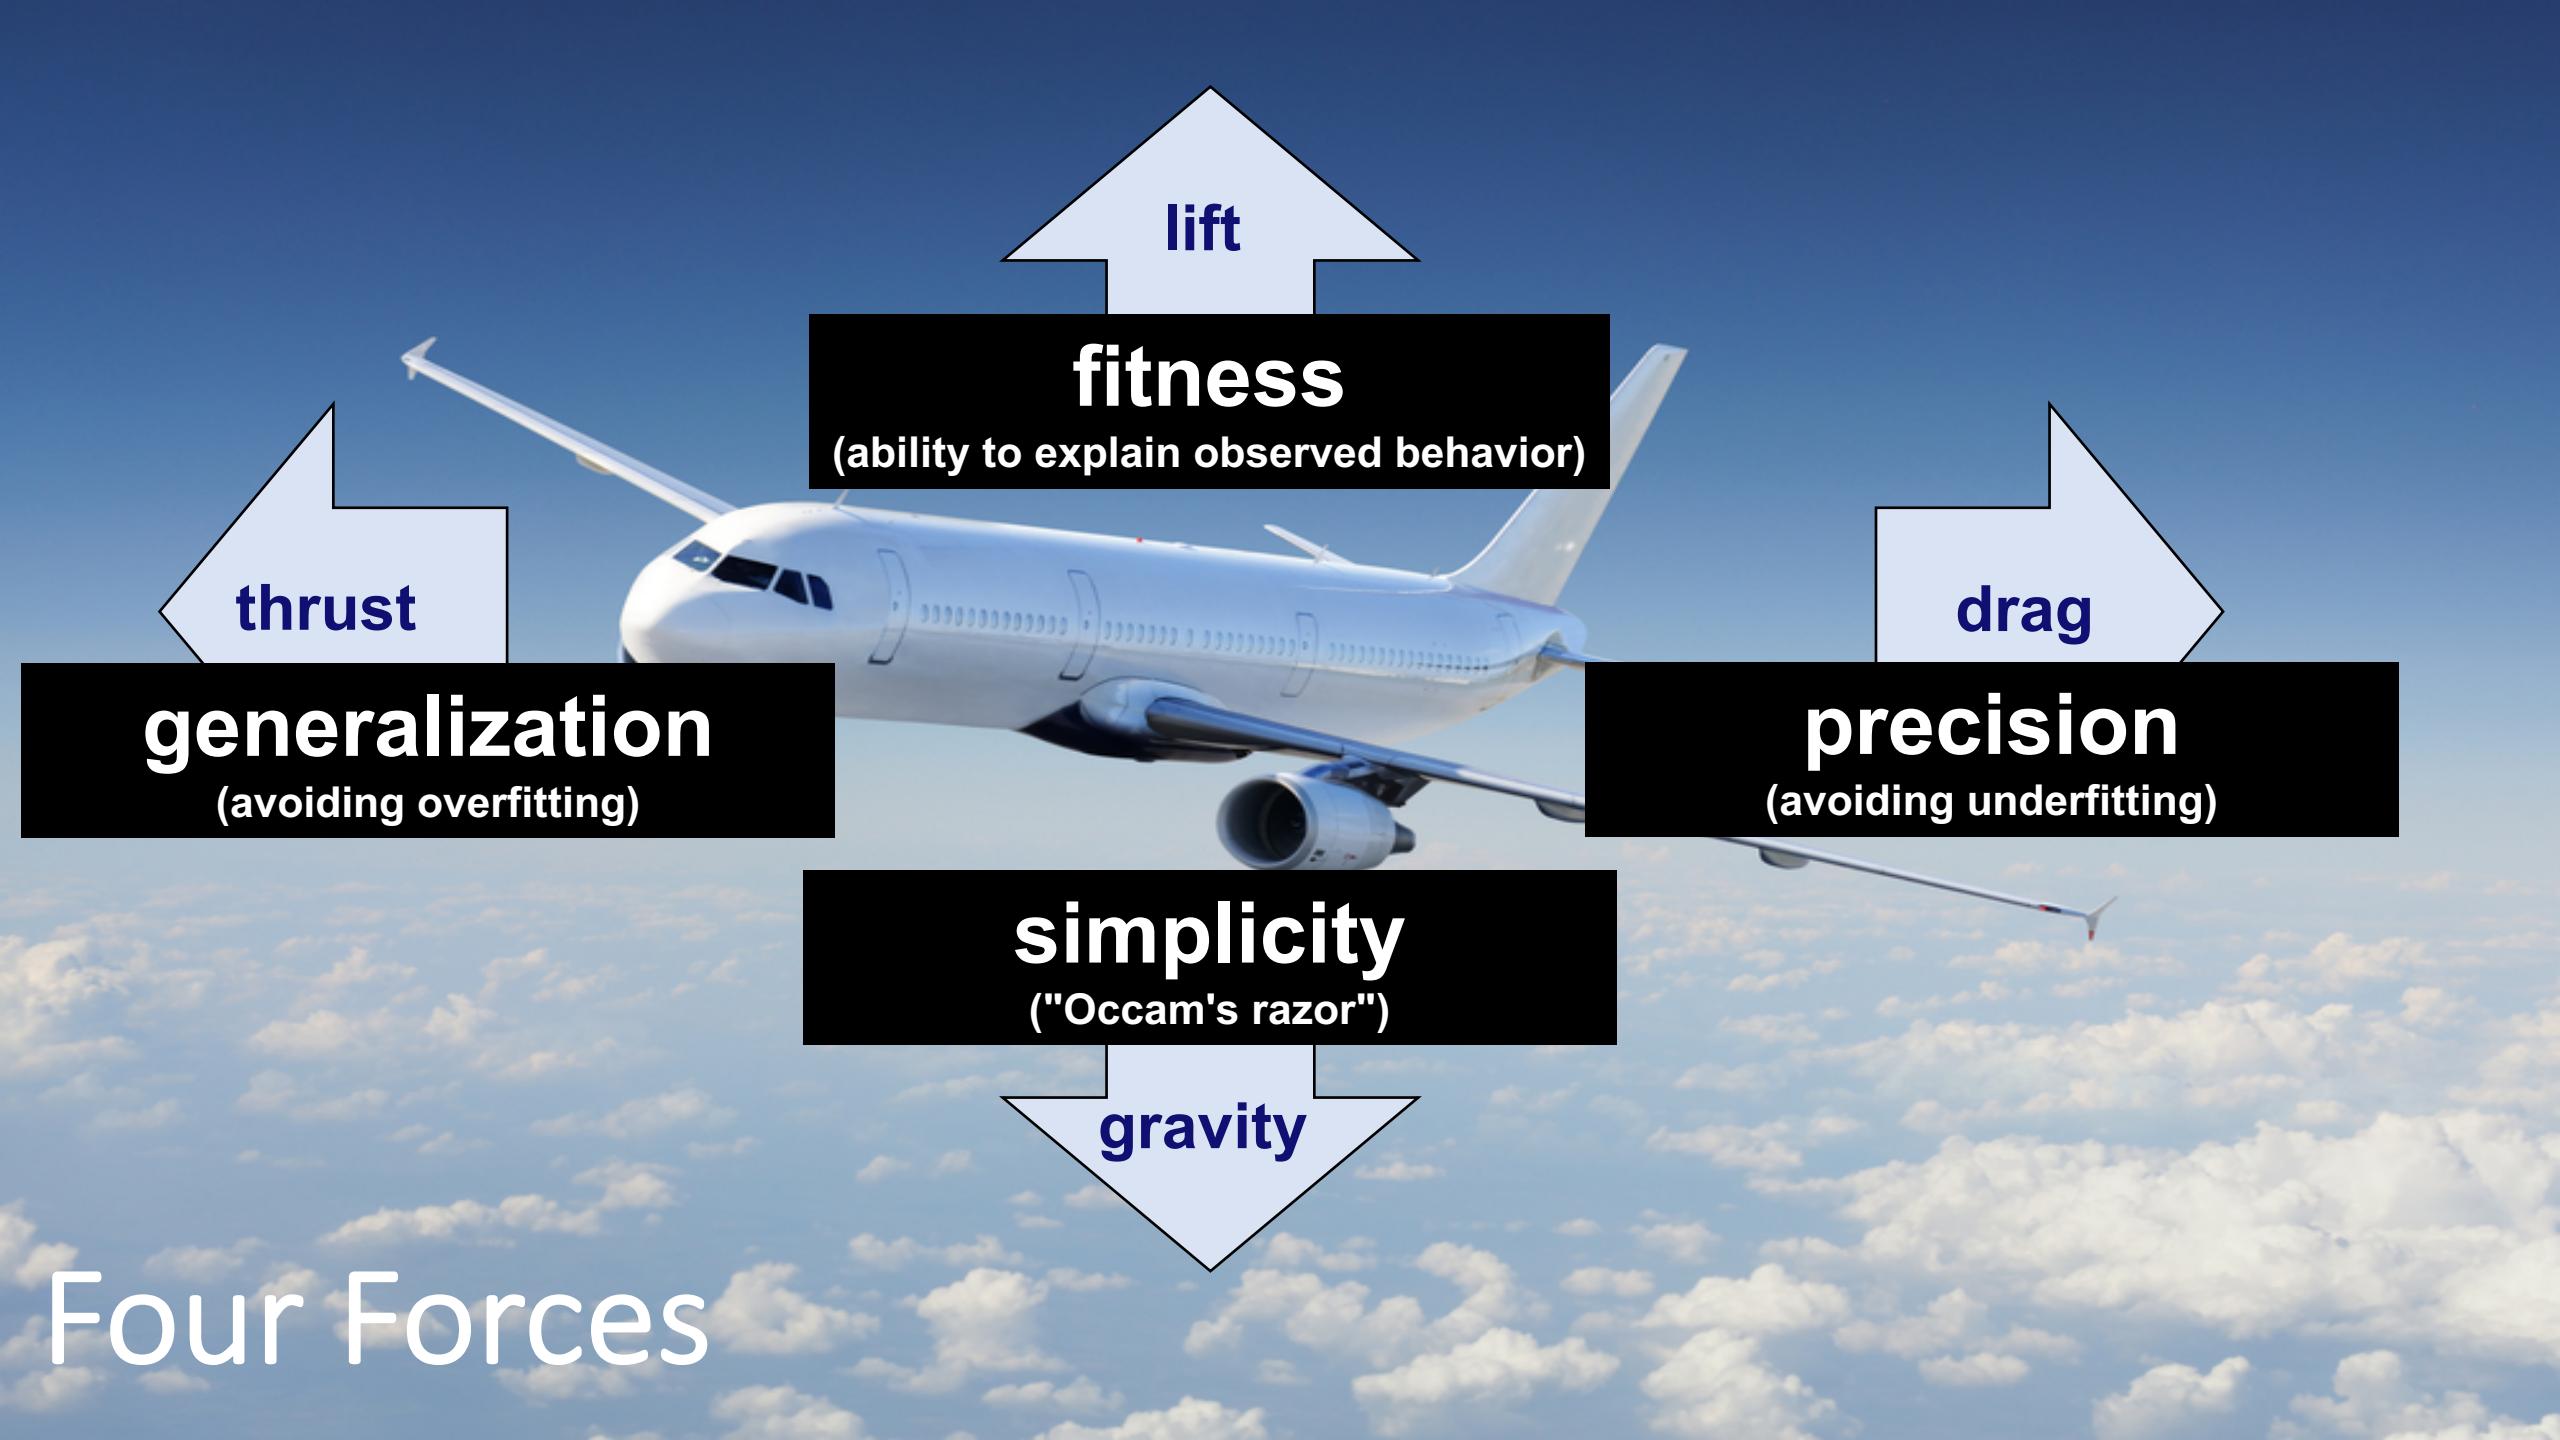
The Spectrum of Conformance Checking

- Conformance checking is about detecting and analyzing deviations between **observed/recorded** and **modelled** behavior.
- Modelled behaviour: Process models
- Observed behaviour: Event data (e.g., event logs)
- **Conformance Artefact Construction:**
 - **Replaying** event data on process models allows to discover deviations.
 - Two replay methods:
 - Token-Replay (intuitive, heuristic)
 - **Alignments** (sound, optimal).

Overview of Behaviours



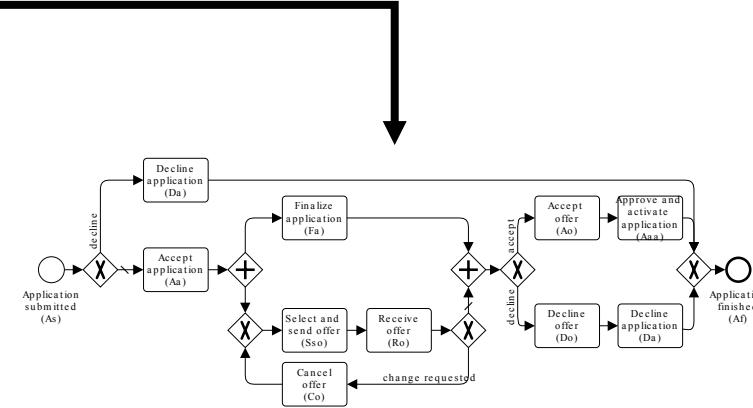
Behaviour as recorded in an event log L, as captured in a process model M, and as shown by a system/process S



Principal Dimensions: fitness and precision

Fitness: Has the recorded behaviour been modelled?

Event	Application	Offer	Activity	Amount	Signed	Timestamp
e ₁₃	A5634	Application submitted	€2,000	Jan 01, 12:31
e ₁₄	A5634	Accept application	€2,000	Jan 01, 12:32
e ₁₅	A5635	Application submitted	€5,000	Jan 02, 04:31
e ₁₆	A5635	Accept application	€5,000	Jan 02, 04:32
e ₁₇	A5636	Application submitted	€200	Jan 03, 06:59
e ₁₈	A5636	Accept application	€200	Jan 03, 07:00
....
e ₂₂	A5634	Finalise application	Jan 03, 09:00
e ₂₃	A5636	Finalise application	Jan 03, 09:01
e ₂₄	A5635	Decline application	Jan 03, 09:02
e ₂₅	A5635	Decline application	Jan 03, 09:03
....
e ₃₀	A5636	03521	Select and send offer	€500	Jun 04, 16:32
....
e ₃₇	A5634	03541	Select and send offer	€1,500	Jun 05, 12:22
e ₃₈	A5636	03521	Receive offer	NO	Jun 05, 12:33
e ₃₉	A5636	03521	Cancel offer	NO	Jun 05, 12:34
e ₄₀	A5636	03542	Select and send offer	€500	Jun 05, 13:29
e ₄₁	A5636	03542	Receive offer	YES	Jun 08, 08:33
e ₄₂	A5634	03542	Accept offer	NO	Jun 08, 16:34
e ₄₃	A5634	03541	Receive offer	NO	Jun 10, 10:00
....
e ₅₄	A5634	03541	Decline offer	Jun 10, 10:04
....
e ₆₄	A5634	Decline application	Jan 10, 10:05
e ₆₅	A5634	Application finished	Jan 10, 10:06
e ₆₆	A5636	Approve and activate application	Jan 10, 10:07
e ₆₇	A5636	Application finished	Jan 10, 10:08



Precision: Has the modelled behaviour been recorded ?

Evaluation of Mined Process Models

Motivation

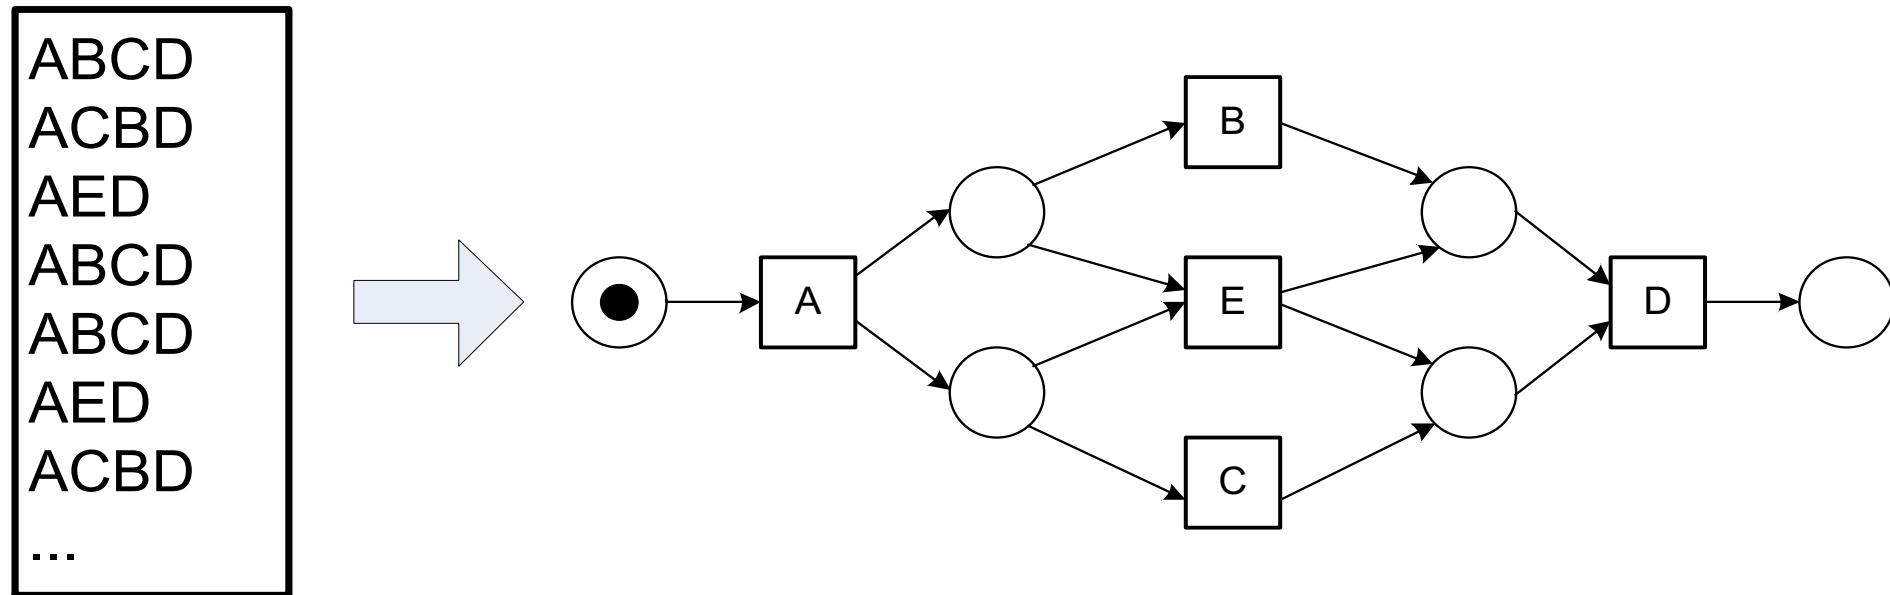
- Different process mining algorithms create different models
- How to judge quality of the created models?

Evaluation of mined process models

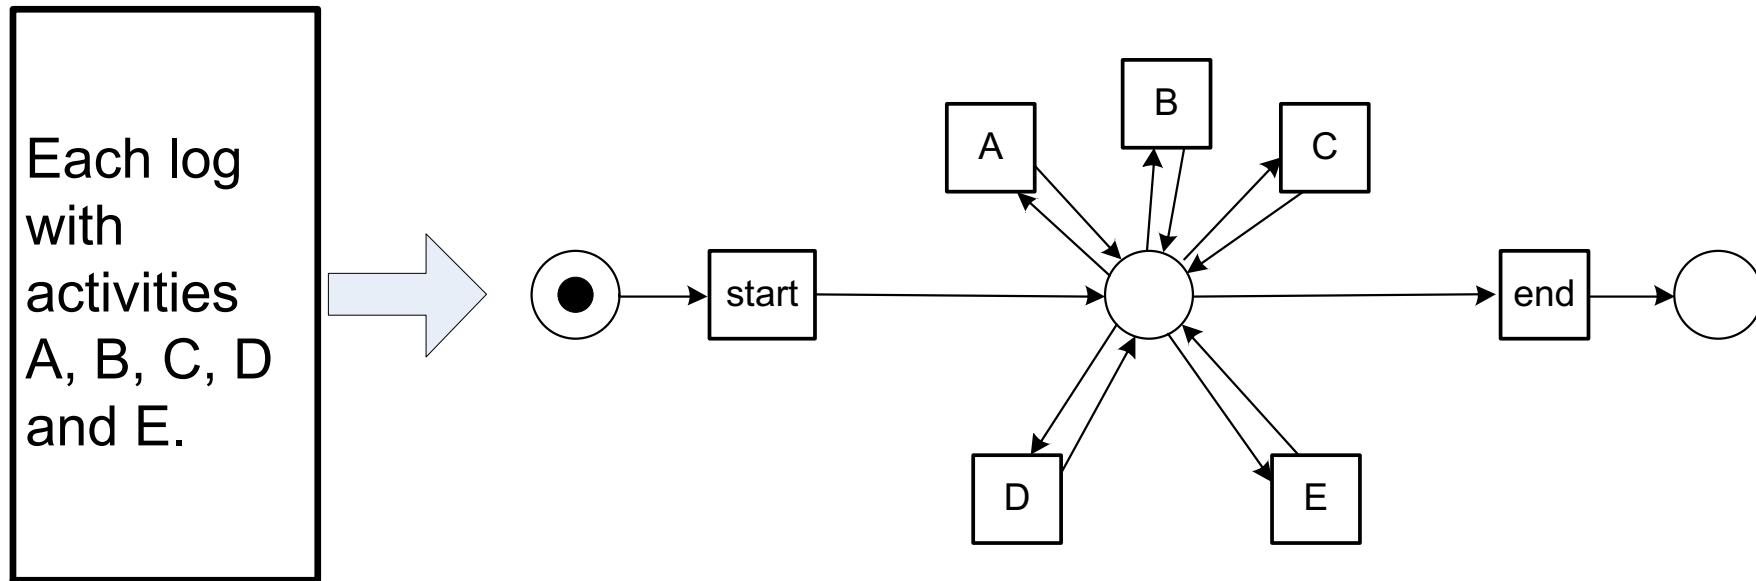
- To compare process mining algorithms
- To make pros and cons of different approaches explicit

Evaluation in terms of overfitting/underfitting

Overfitting/Underfitting



Overfitting/Underfitting cont.



Evaluation Dimensions

Fitness

- Is the observed behaviour covered by the model?

The Precision – Generalisation Trade-off

- Recall: in general, it is extremely unlikely that a data sample such as a log contains all (!) possible behaviour
 - Does the model allow for only the observed behaviour?
 - Does the model allow for more than the observed behaviour?

Simplicity

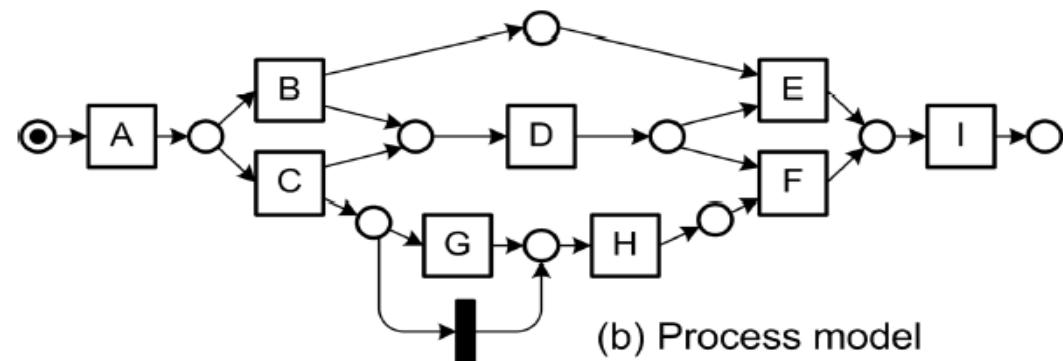
- Does the model have the minimal structure to represent the behaviour?

Example Process Model

- High fitness
- High precision

No. of Instances	Log Traces
1207	ABDEI
145	ACDGHFI
56	ACGDHFI
23	ACHDFI
28	ACDHFI

(a) Event Log



(b) Process model

“Sequence”

- Low fitness
- High precision

No. of Instances	Log Traces
1207	ABDEI
145	ACDGHFI
56	ACGDHFI
23	ACHDFI
28	ACDHFI

(a) Event Log



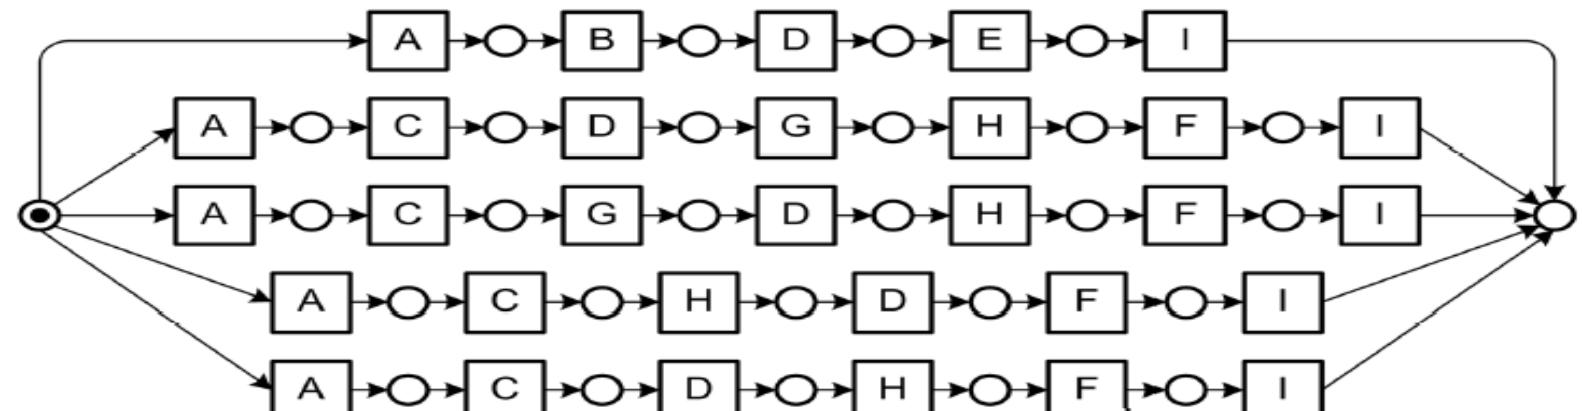
(c) Process model

“Log”

- High fitness
- High precision
- No generalisation
- Bad structure

No. of Instances	Log Traces
1207	ABDEI
145	ACDHFI
56	ACGDHFI
23	ACHDFI
28	ACDHFI

(a) Event Log



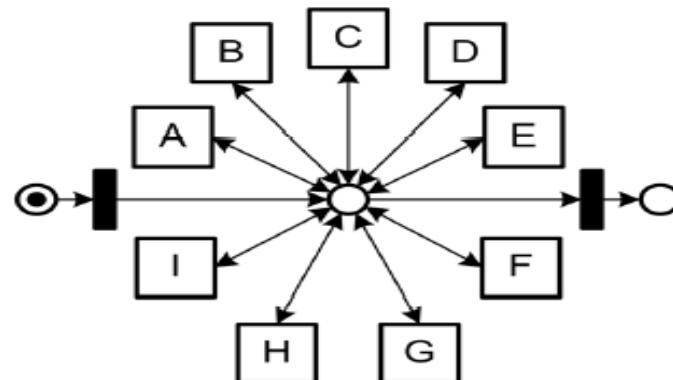
(e) Process model

“Flower”

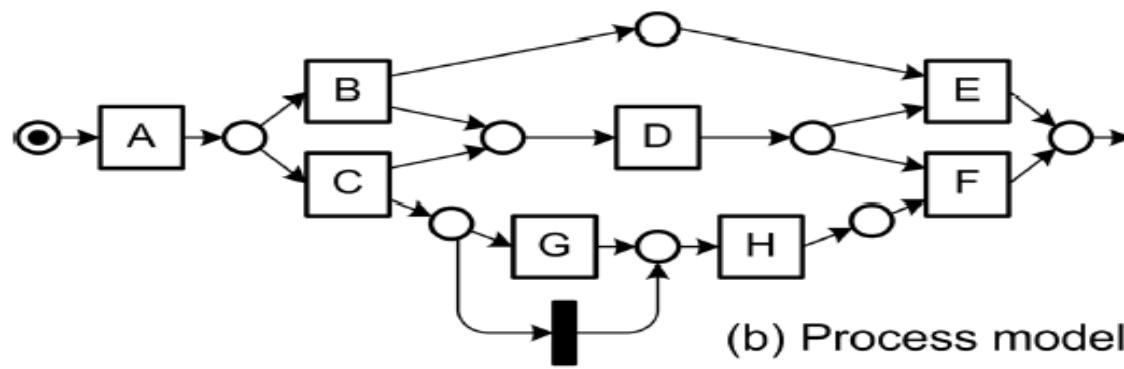
- High fitness
- Low precision

No. of Instances	Log Traces
1207	ABDEI
145	ACDGHFI
56	ACGDHFI
23	ACHDFI
28	ACDHFI

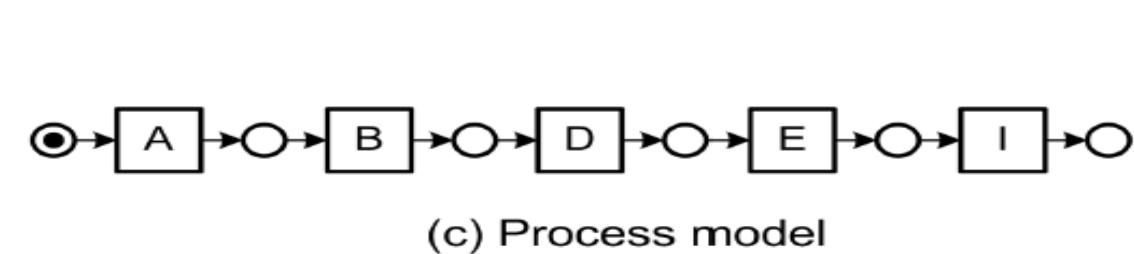
(a) Event Log



(d) Process model



(b) Process model



(c) Process model

*fitness +
precision +
generalization +
structure +*

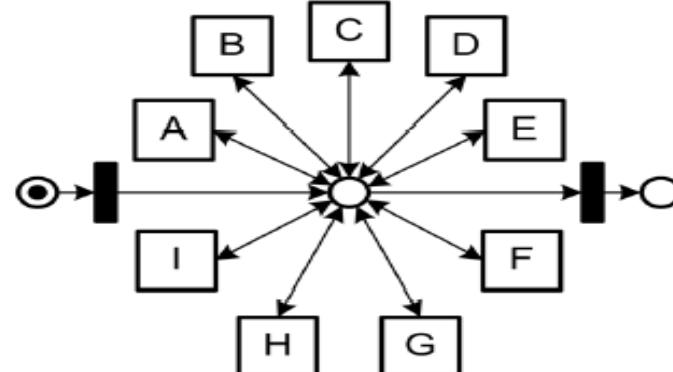
*fitness +
precision -
generalization +
structure +*

*fitness -
precision +
generalization -
structure +*

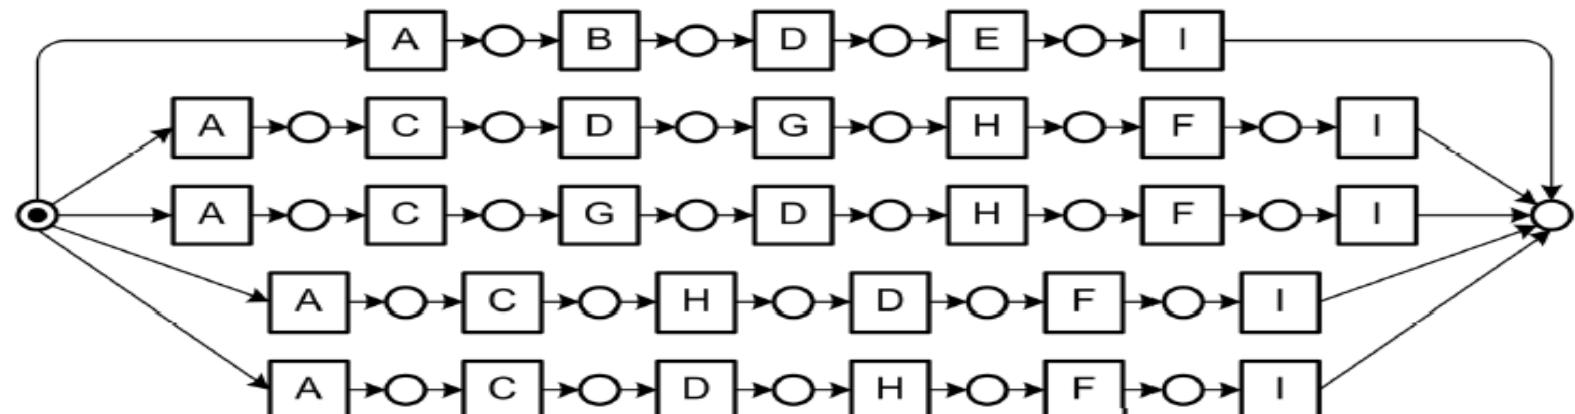
*fitness +
precision +
generalization -
structure -*

No. of Instances	Log Traces
1207	ABDEI
145	ACDGHFI
56	ACGDHFI
23	ACHDFI
28	ACDHFI

(a) Event Log



(d) Process model



(e) Process model

Fitness Measure

Given a process model and a process log

- How to quantify the behaviour represented by the traces w.r.t. the model?

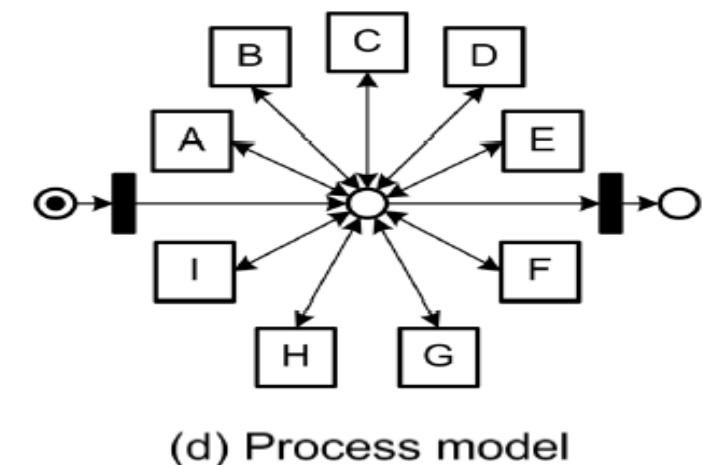
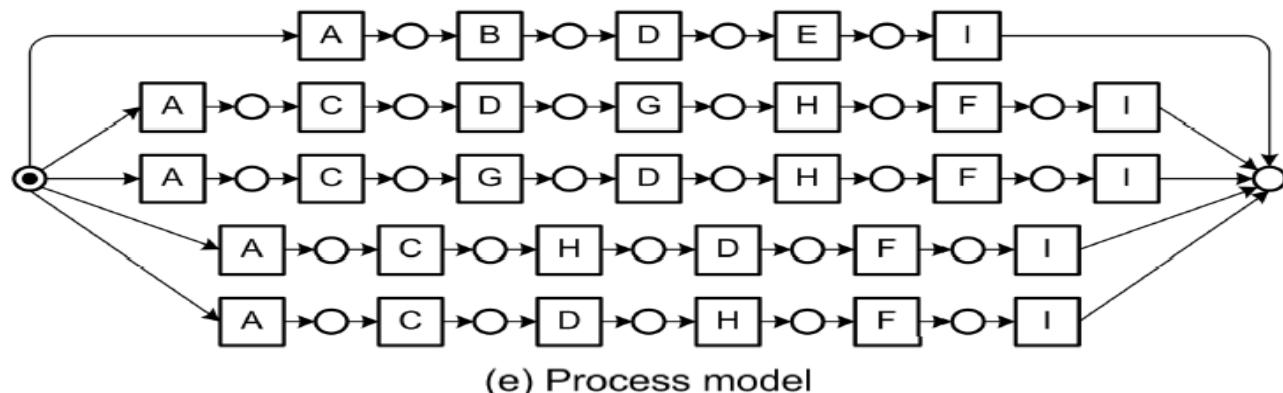
Idea

- Try to “replay” each trace in the model
- Try to fire transitions according to the trace
- Analyse the relation between log and model semantics
- Not ok:
 - Transition fires although it is not enabled
 - Tokens remain in the net and are not consumed

Precision Measures

Given a process model and a process log

- How to quantify the relation between “*observed behaviour / possible behaviour*”?
- How much of the behaviour of the model is observed in the log?
- Or, in terms of generalisation: how much behaviour of the model is not observed in the log?



Simplicity

Goal: find a *simple* model

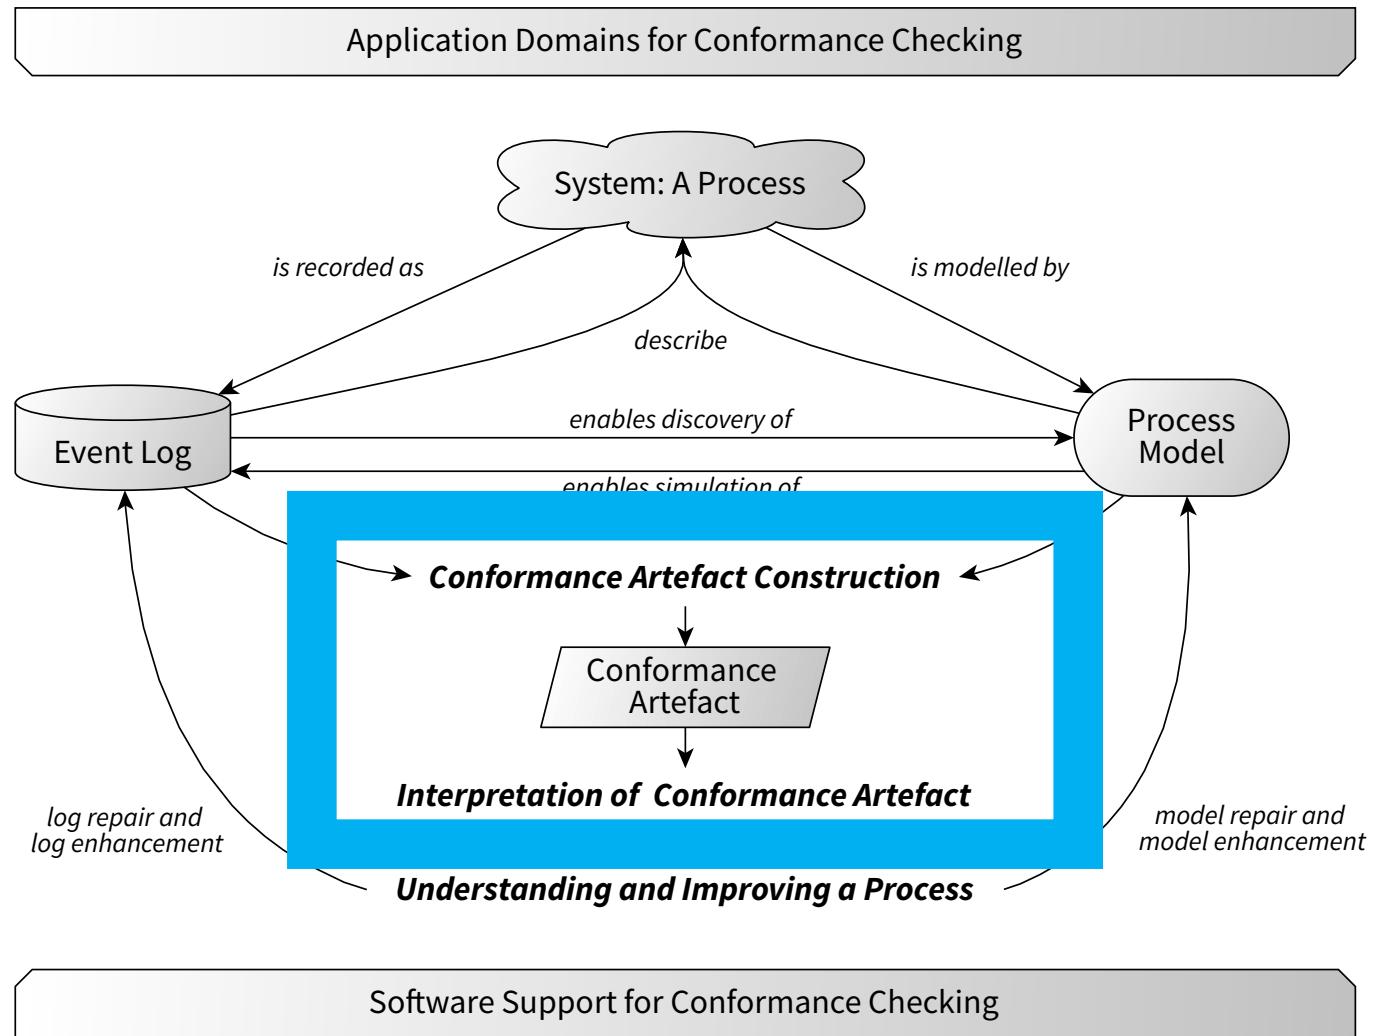
- ... that meets the requirements in terms of fitness and precision/generalisation
- ... to chose among a set of candidate models

Yet, different notions of simple are meaningful:

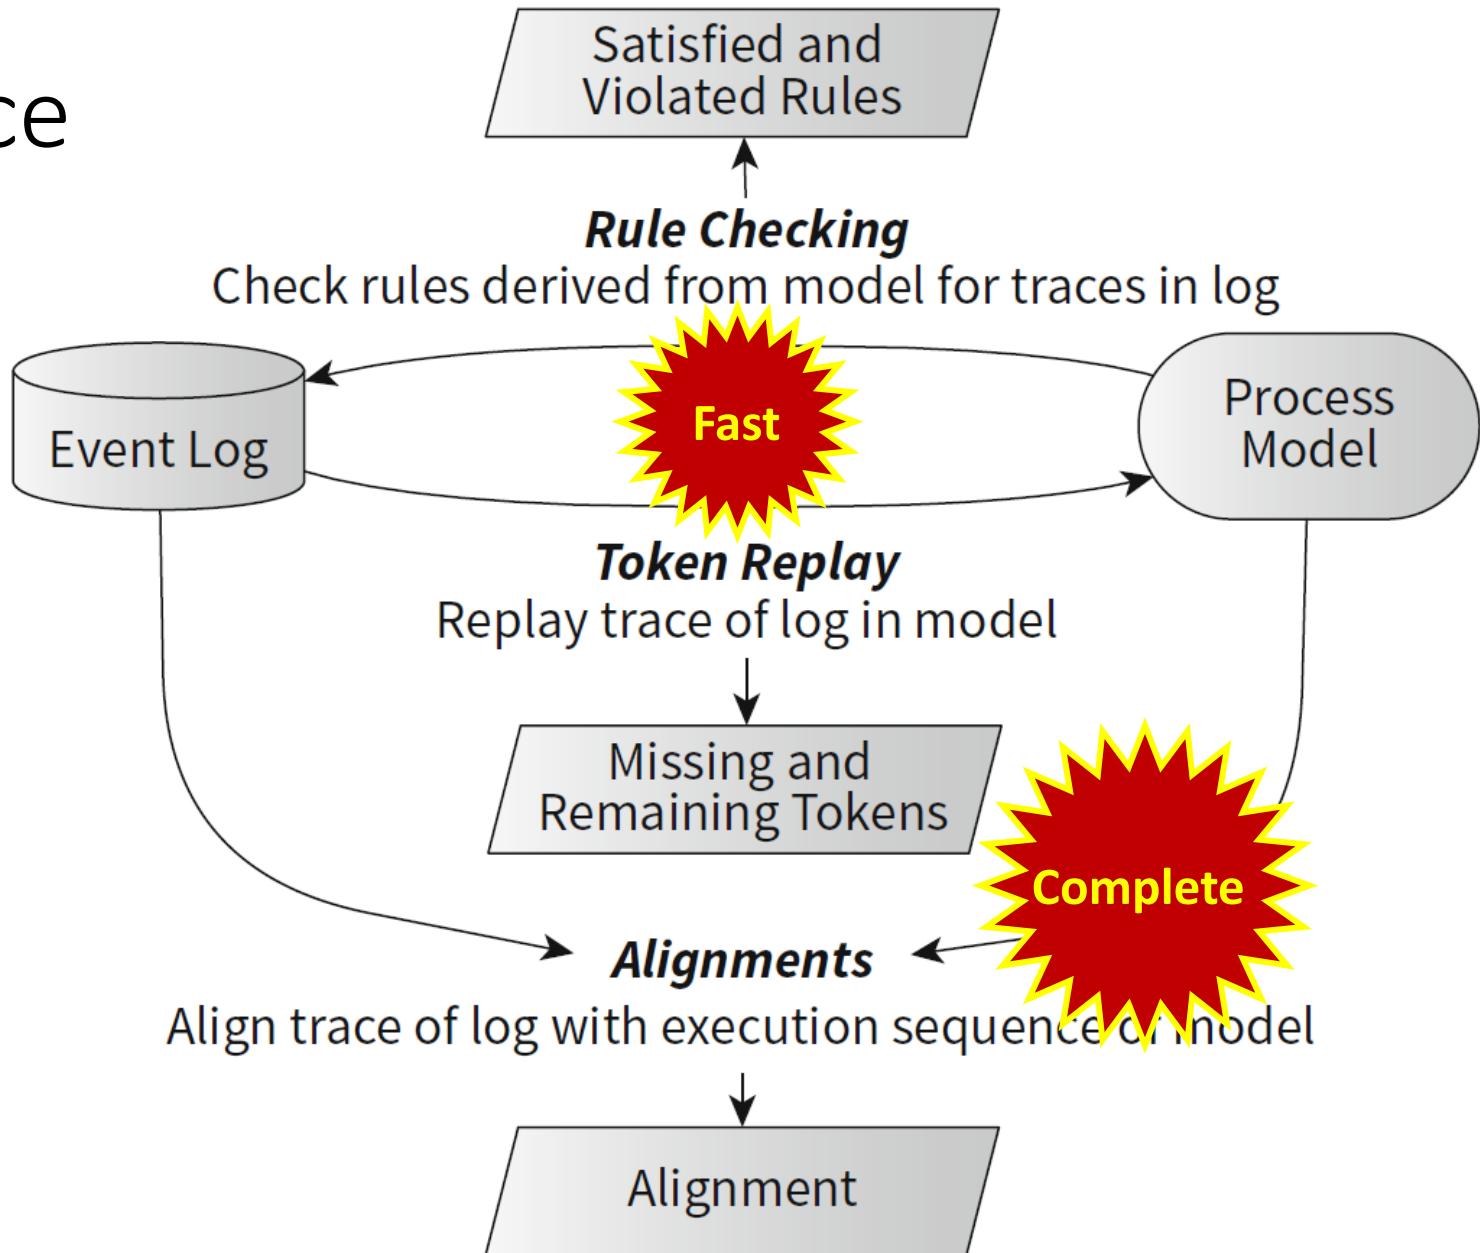
- Depends on modelling language
- Depends on understandability considerations
- Mostly related to model structure, yet may also concern semantics

Baseline measure: Simplicity as the number of model elements

The Spectrum of Conformance Checking



Computing Conformance Artefacts



Computing Conformance Checking Artefacts

Rule Checking

Idea of Rule Checking

General idea:

- Process model constrains the possible behaviour of a process
- Formalise constraints by rules that execution sequences have to obey to
- Then, verify these rules for the traces of a log

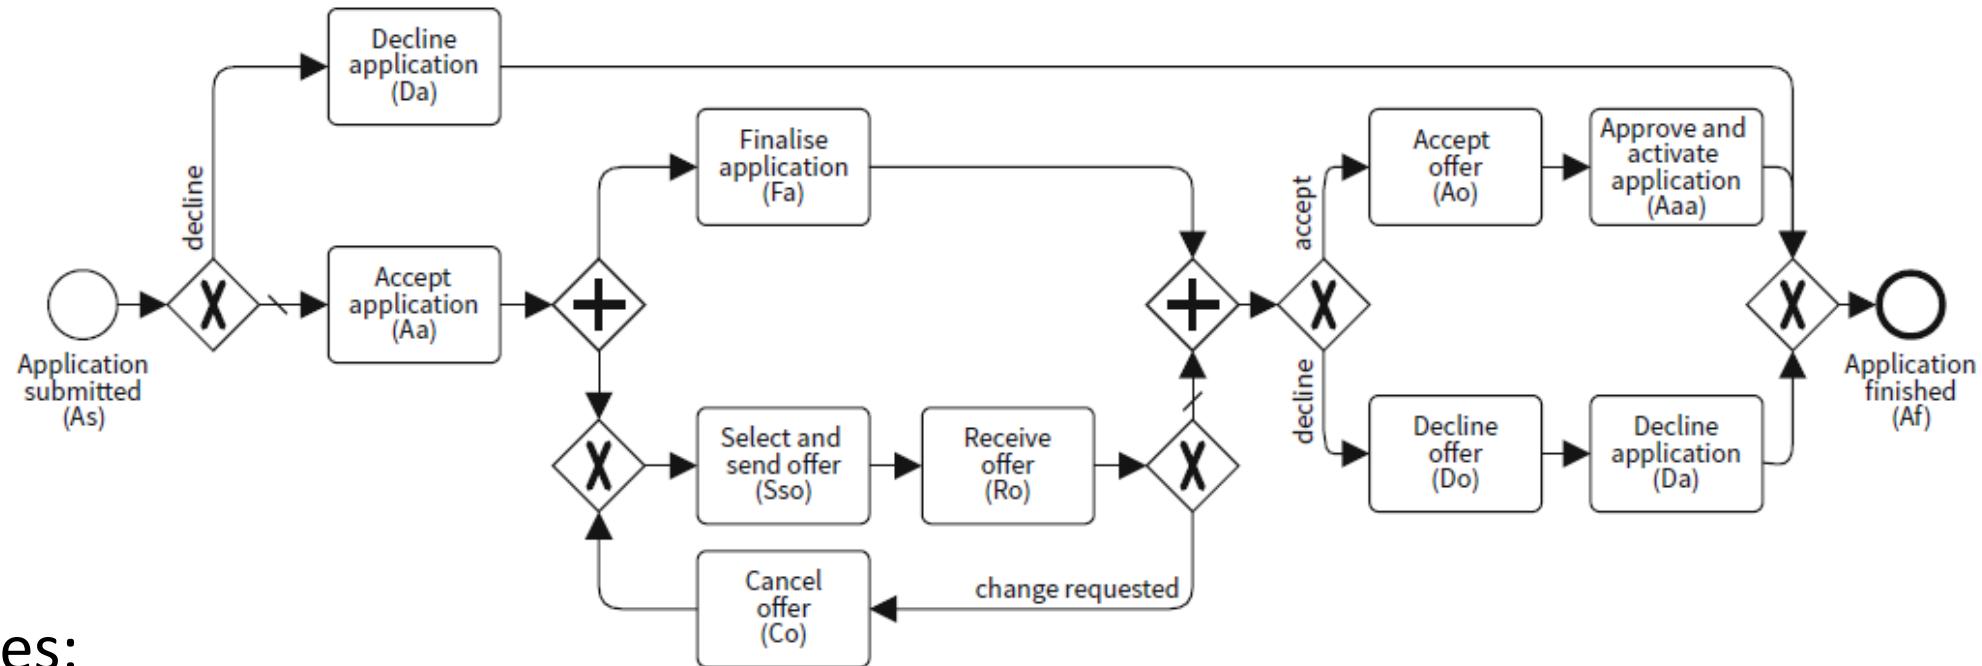
Rule checking focuses on the fitness dimension

- Traces are fitting, if they satisfy the rules
- Traces are non-fitting, if rules are violated

The dimension of precision is not targeted

- Specificity of rules is neglected
- Many more traces, not contained in the log, could also satisfy the rules

Illustration of Rules



Rules:

- R1: An application can be accepted (Aa) at most once
- R2: An accepted application (Aa), that must have been submitted (As) earlier, and eventually an offer needs to be selected and sent (Sso) for it
- R3: An application must never be finalised (Fa), if the respective offer has been declined (Do) already
- R4: An offer is either accepted (Ao) or declined (Do), but cannot be both accepted and declined

What rules?

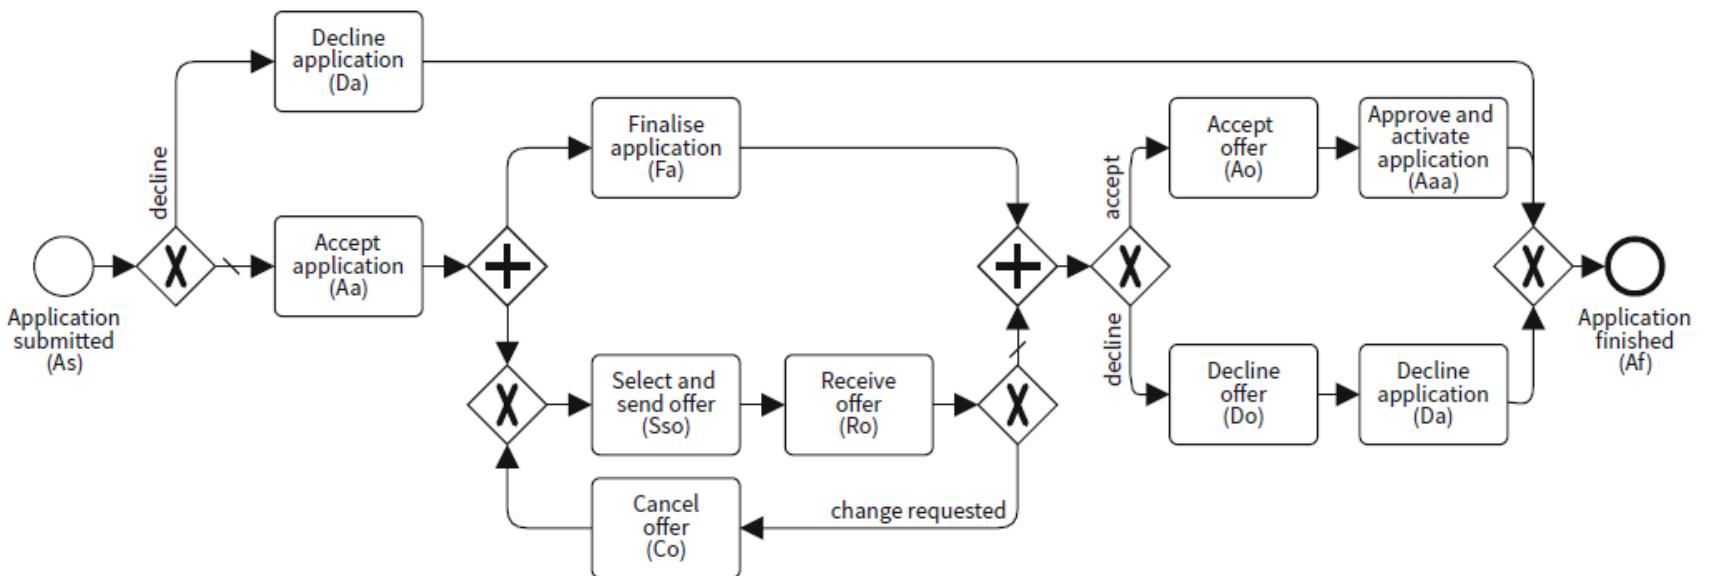
Often: unary and binary rules over activities

- Specify how a single activity is executed or how the executions of a pair of activities relate to each other
- N-ary rules would lead to combinatorial explosion

Common types of rules:

- Cardinality rules: Upper and/or lower bound for the number of executions of an activity
- Precedence and response rules: An activity is always preceded/succeeded by another one
- Ordering rules: Activities are independent, but if they are both executed, they are ordered
- Exclusiveness rules: Activities must not be executed in the same case

Cardinality Rules



Activity	<i>As</i>	<i>Da</i>	<i>Aa</i>	<i>Fa</i>	<i>Sso</i>	<i>Ro</i>	<i>Co</i>	<i>Ao</i>	<i>Aaa</i>	<i>Do</i>	<i>Af</i>
----------	-----------	-----------	-----------	-----------	------------	-----------	-----------	-----------	------------	-----------	-----------

Cardinality:	[1, 1]	[0, 1]	[0, 1]	[0, 1]	[0, n]	[0, n]	[0, n]	[0, 1]	[0, 1]	[0, 1]	[1, 1]
--------------	--------	--------	--------	--------	--------	--------	--------	--------	--------	--------	--------

$T_1 = \langle As, Aa, Sso, Ro, Ao, Aaa, Aaa \rangle$

✓	✓	✓	✓	✓	✓	✓	✓	✓	X	✓	X
---	---	---	---	---	---	---	---	---	---	---	---

$T_2 = \langle As, Sso, Fa, Ro, Co, Ro, Aaa, Af \rangle$

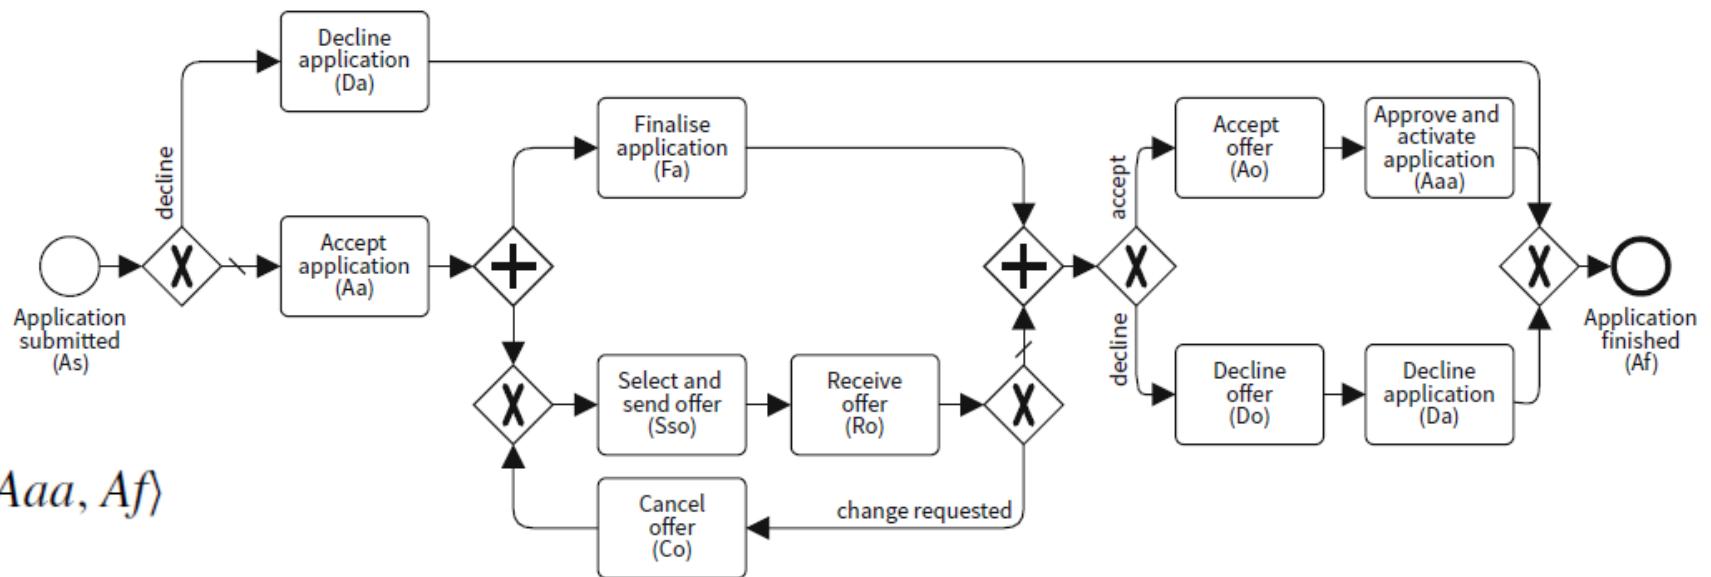
✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
---	---	---	---	---	---	---	---	---	---	---	---

$T_3 = \langle As, Aa, Sso, Ro, Fa, Ao, Do, Da, Af \rangle$

✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
---	---	---	---	---	---	---	---	---	---	---	---

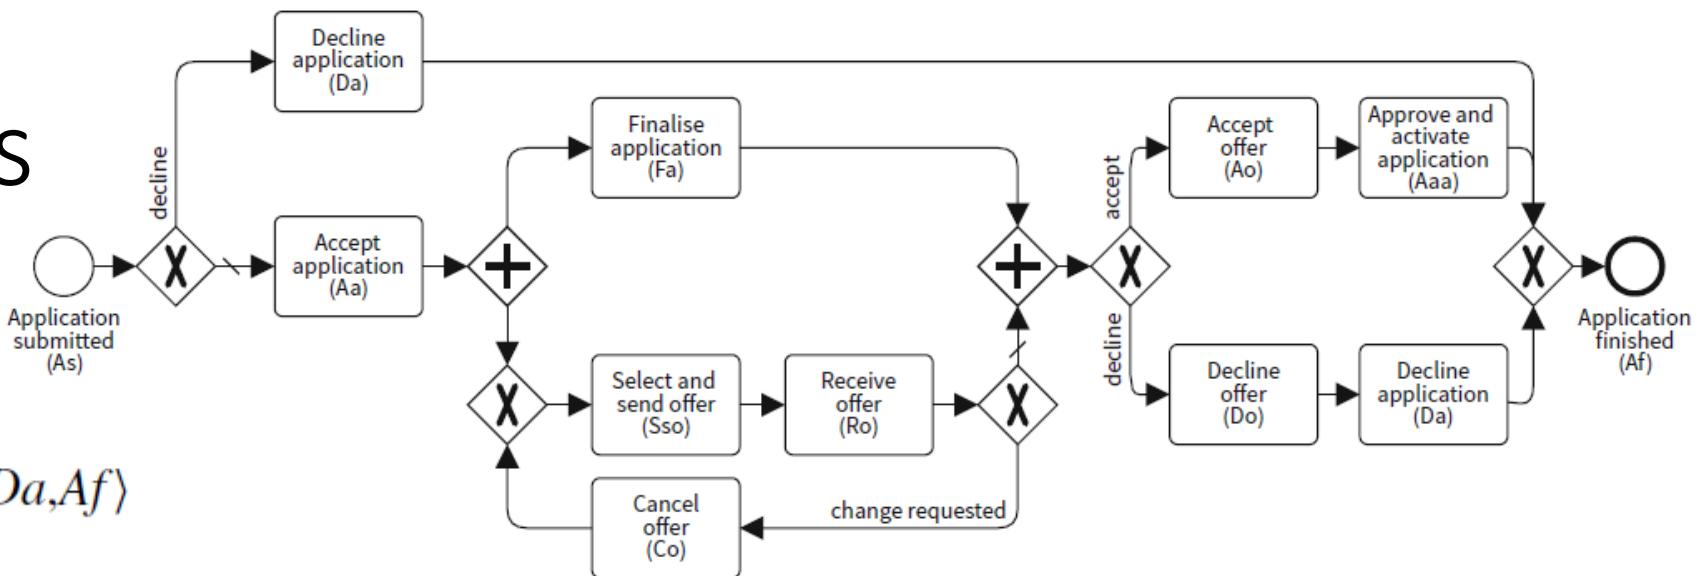
Precedence Rules

$$T_2 = \langle As, Sso, Fa, Ro, Co, Ro, Aaa, Af \rangle$$



	<i>As</i>	<i>Da</i>	<i>Aa</i>	<i>Fa</i>	<i>Sso</i>	<i>Ro</i>	<i>Co</i>	<i>Ao</i>	<i>Aaa</i>	<i>Do</i>	<i>Af</i>
<i>As</i>											
<i>Da</i>	✓										
<i>Aa</i>	✓										
<i>Fa</i>	✓				X						
<i>Sso</i>	✓				X						
<i>Ro</i>	✓				X			✓			
<i>Co</i>	✓				X		✓	✓			
<i>Ao</i>	✓				✓	✓	✓	✓			
<i>Aaa</i>	✓				X	✓	✓	✓			X
<i>Do</i>	✓				✓	✓	✓	✓			
<i>Af</i>	✓										

Exclusiveness Rules



$$T_3 = \langle As, Aa, Sso, Ro, Fa, Ao, Do, Da, Af \rangle$$

	As	Da	Aa	Fa	Sso	Ro	Co	Ao	Aaa	Do	Af
As	✓										
Da		✓									
Aa			✓								
Fa				✓							
Sso											
Ro											
Co											
Ao		✗						✓		✗	
Aaa		✓							✓	✓	
Do								✗	✓	✓	
Af											✓

Feedback on Non-Conformance

Rule-checking gives information on fitness of trace

- Is the model well-suited to explain the recorded trace?
- Yes, if rules are satisfied; no, if rules are violated
- Fine-granular feedback on the level of activities

Aggregated feedback on the level of a log

- Rules that are frequently violated point to general conformance issues
- Indicator for hot-spots
- Association rule mining to highlight dependencies between two violations v and v' : Ratio of number of traces showing both violations and number of traces showing only violation v

Fitness Measure

Counting of satisfied rules enables quantification of fitness:

$$\text{fitness}(\sigma, M) = \frac{|\{r \in R_M \mid r \text{ is satisfied by } \sigma\}|}{|R_M|}$$

$$\text{fitness}(L, M) = \frac{|\{r \in R_M \mid r \text{ is satisfied by all } \sigma \in L\}|}{|R_M|}$$

Obviously, the above measures depend on the chosen set of rules

- Values under different rule sets cannot be compared
- Measures are influenced by how many rules are affected by a single conformance issue (e.g., missing execution of an activity)

Computing Conformance Checking Artefacts

Token Replay

Idea of Token Replay

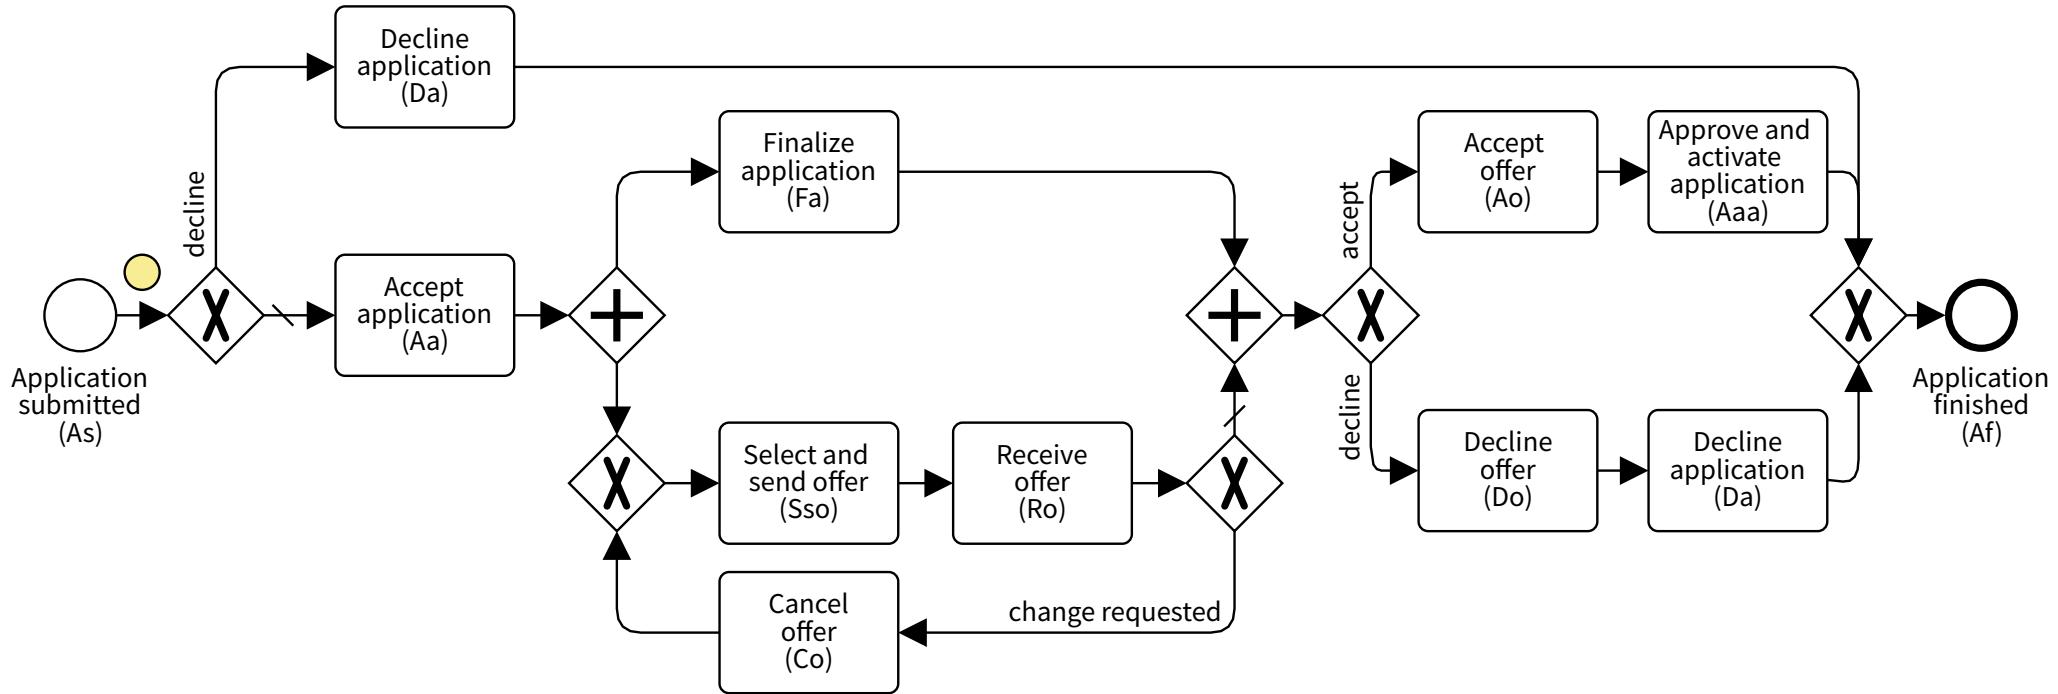
Idea

- Try to “replay” each trace in the model by executing activities according to the trace
- Not ok:
 - Activity is executed although it is not enabled
 - Tokens remain in the model and are not consumed

Procedure:

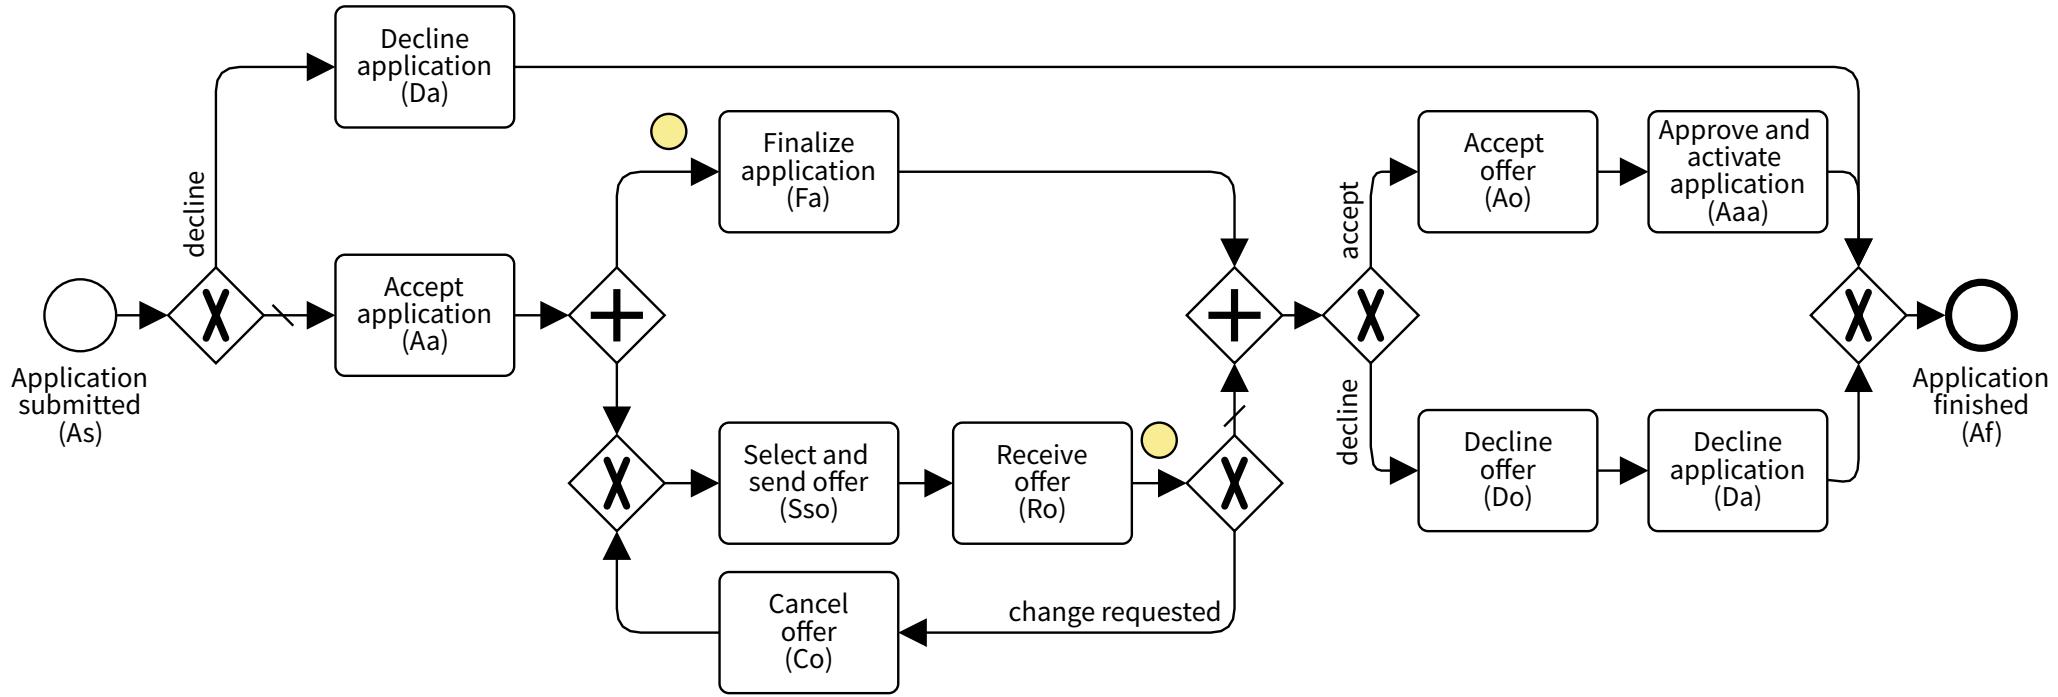
- 1) Initialise the replay procedure by setting the *first event* of the trace as the *current event* of the replay
- 2) Locate the activity for which execution is signalled by the *current event* of the trace in the model
- 3) Replay the *current event* of the trace by executing the identified activity in the *current state* of the model by:
 - Consuming a token on the incoming arc of the activity, even if there is none (if so, record as missing)
 - Producing a token on the outgoing arc of the respective task (if there is any)
- 4) Set the state reached as the *current state* of the process model
- 5) If the current event is not the last in the trace, consider the next event in the trace as the current event and continue the replay from step 2 onwards

Token-Based Replay



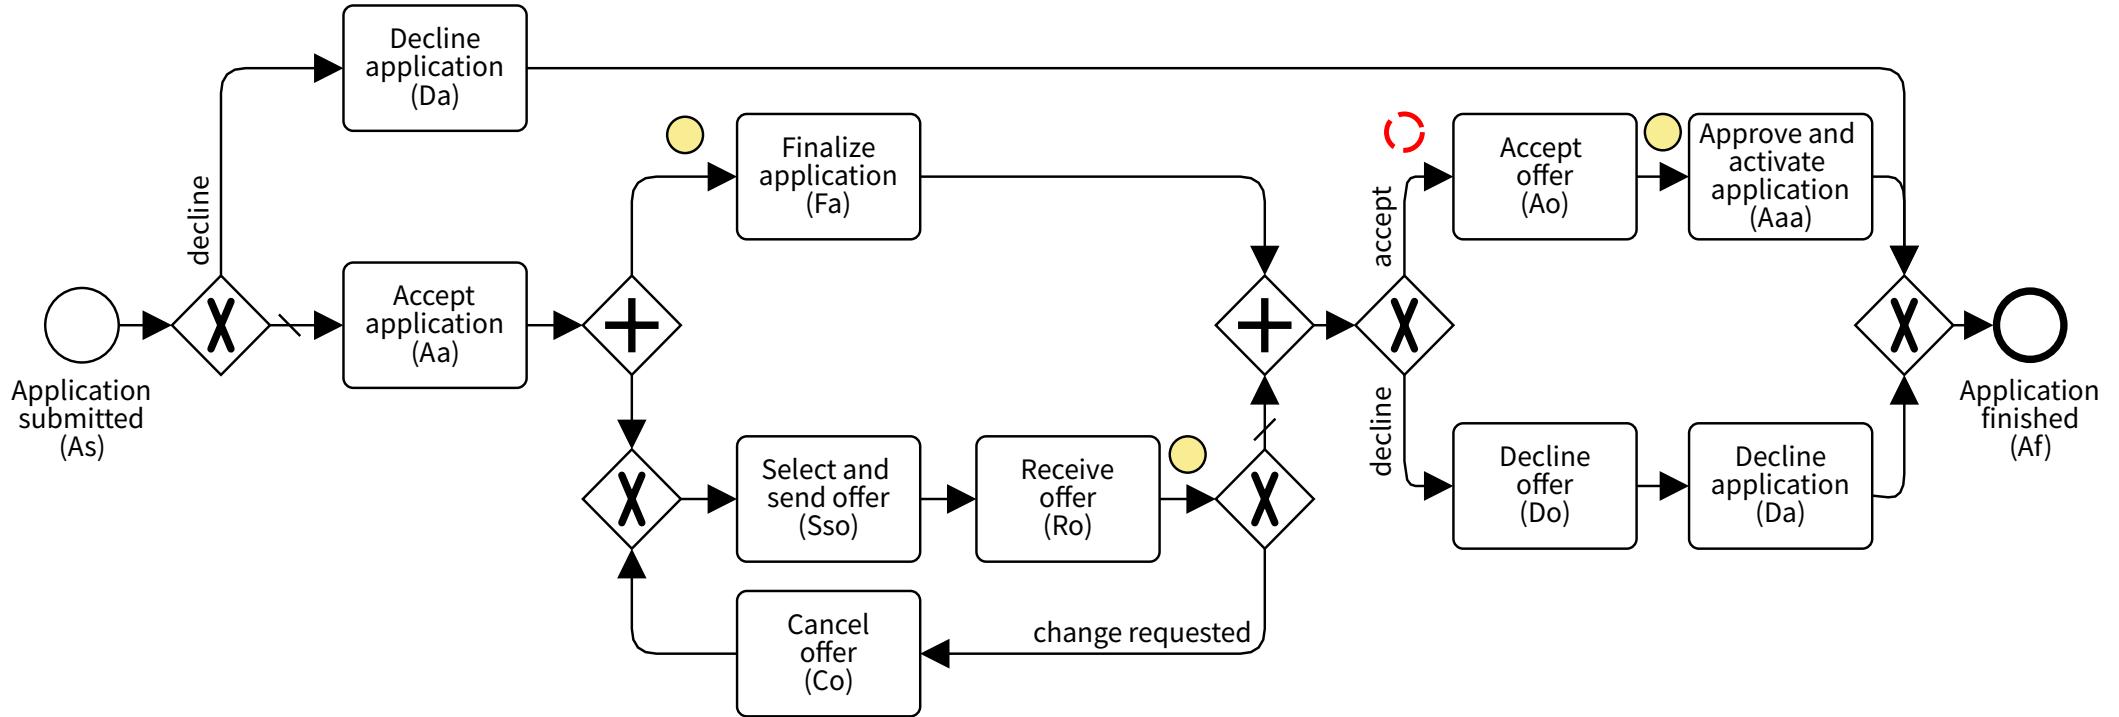
Observed Trace: < As, Aa, Sso, Ro, Ao, Aaa, Aaa >

Token-Based Replay



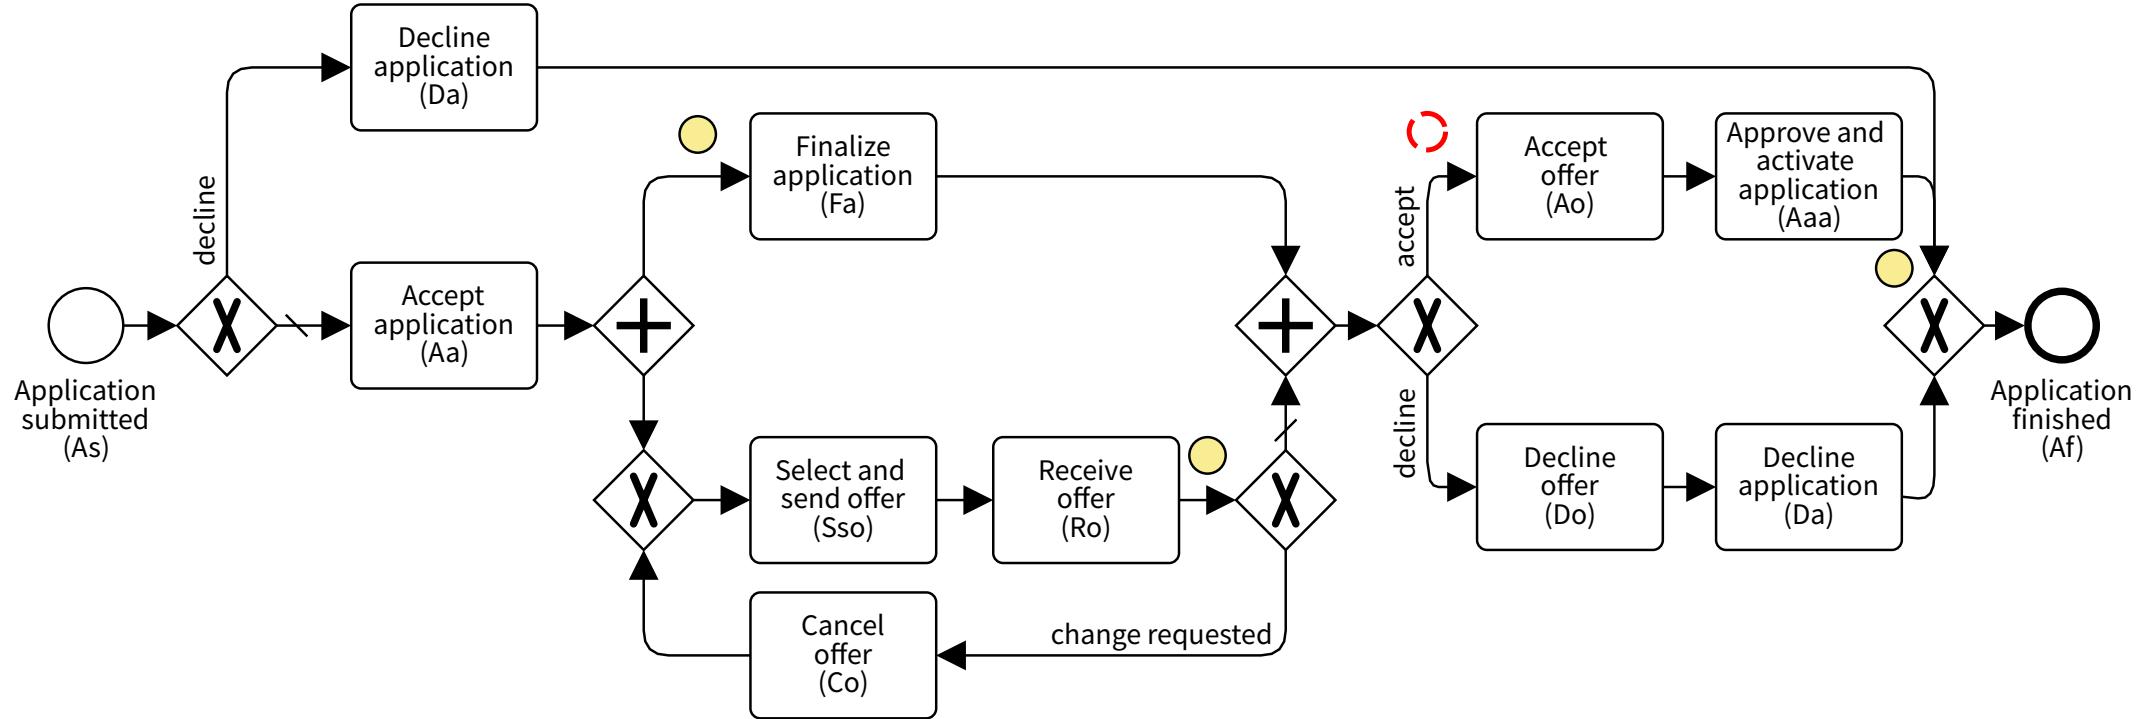
Observed Trace: < As, Aa, Sso, Ro, Ao, Aaa, Aaa >

Token-Based Replay



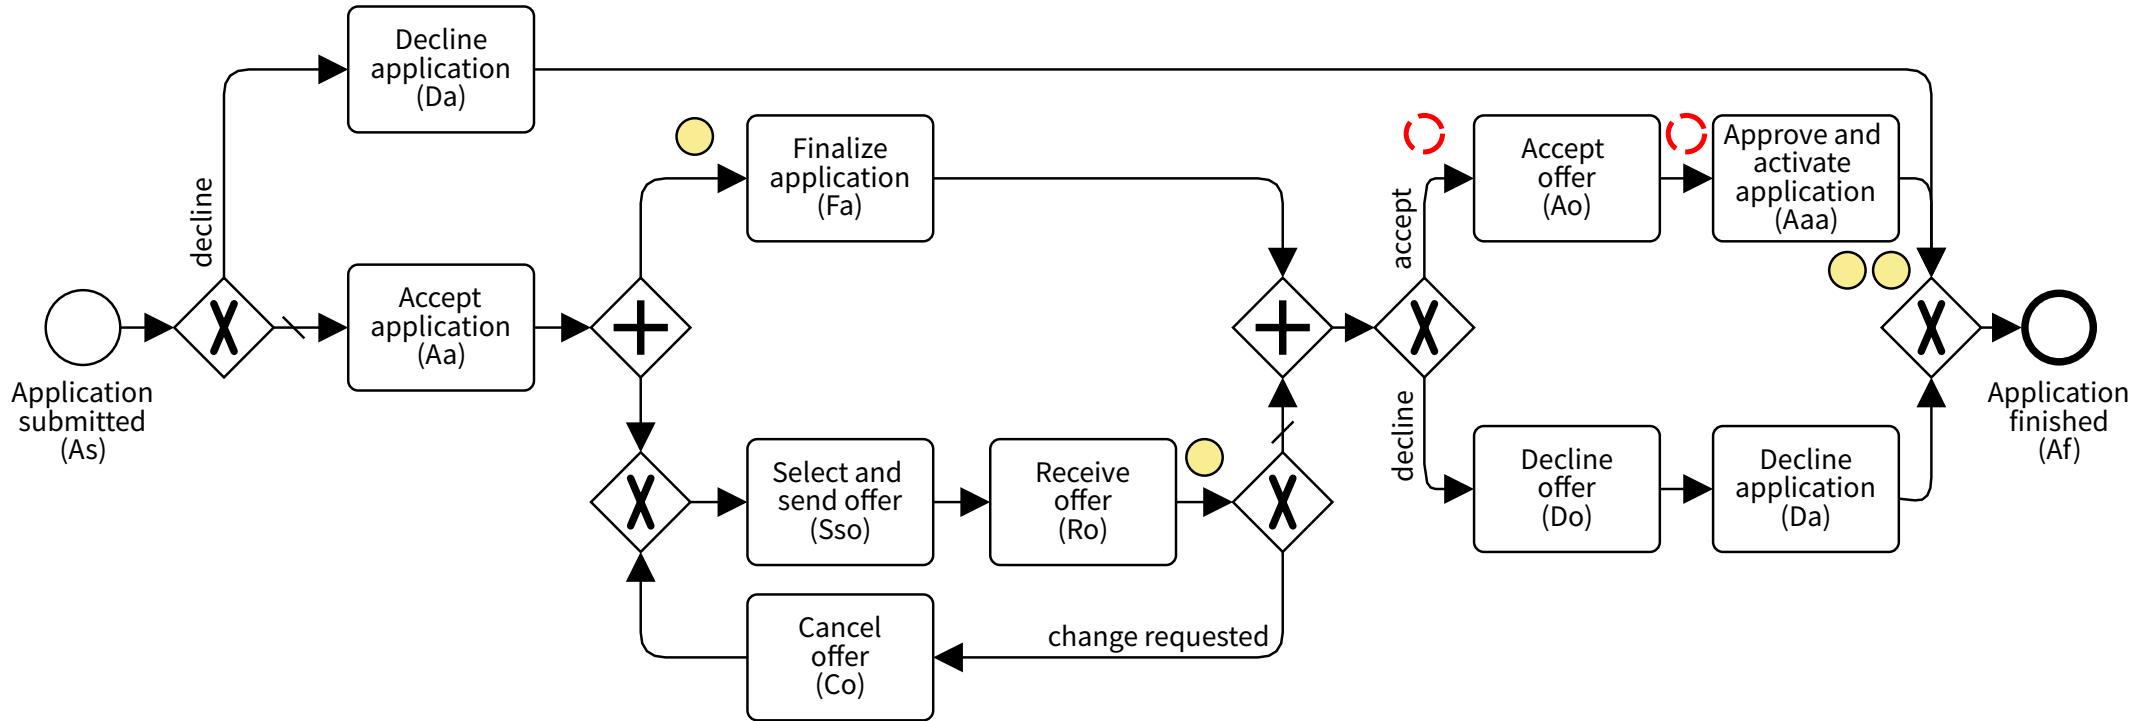
Observed Trace: < As, Aa, Sso, Ro, Ao, Aaa, Aaa >

Token-Based Replay



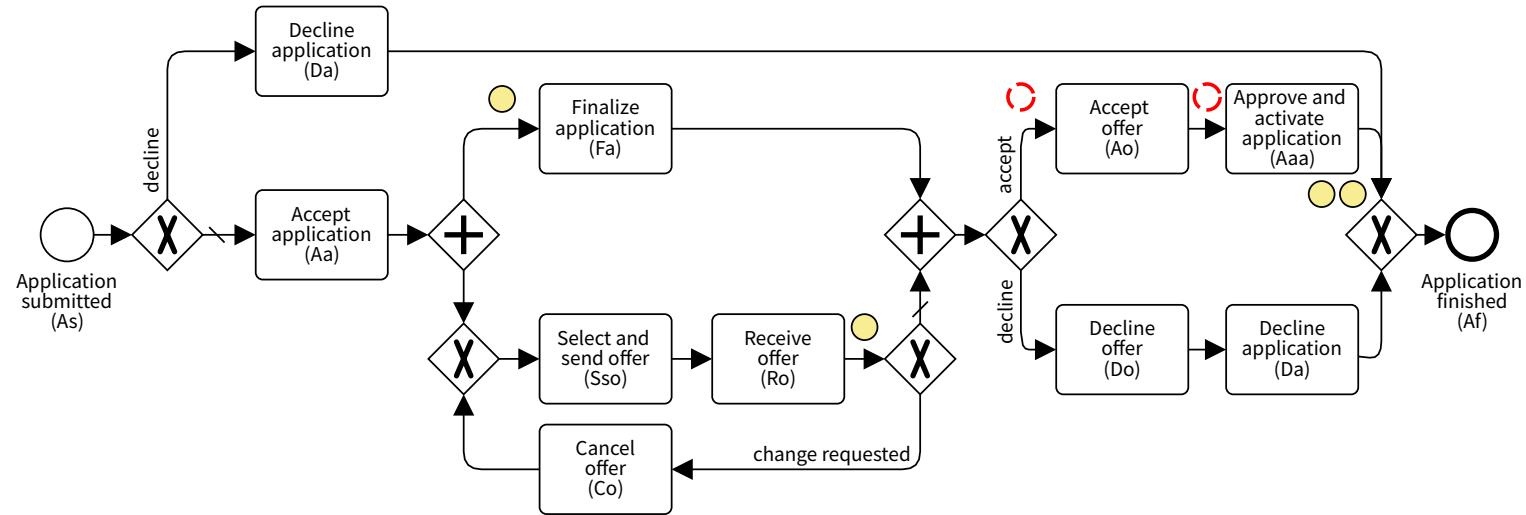
Observed Trace: < As, Aa, Sso, Ro, Ao, Aaa, Aaa >

Token-Based Replay



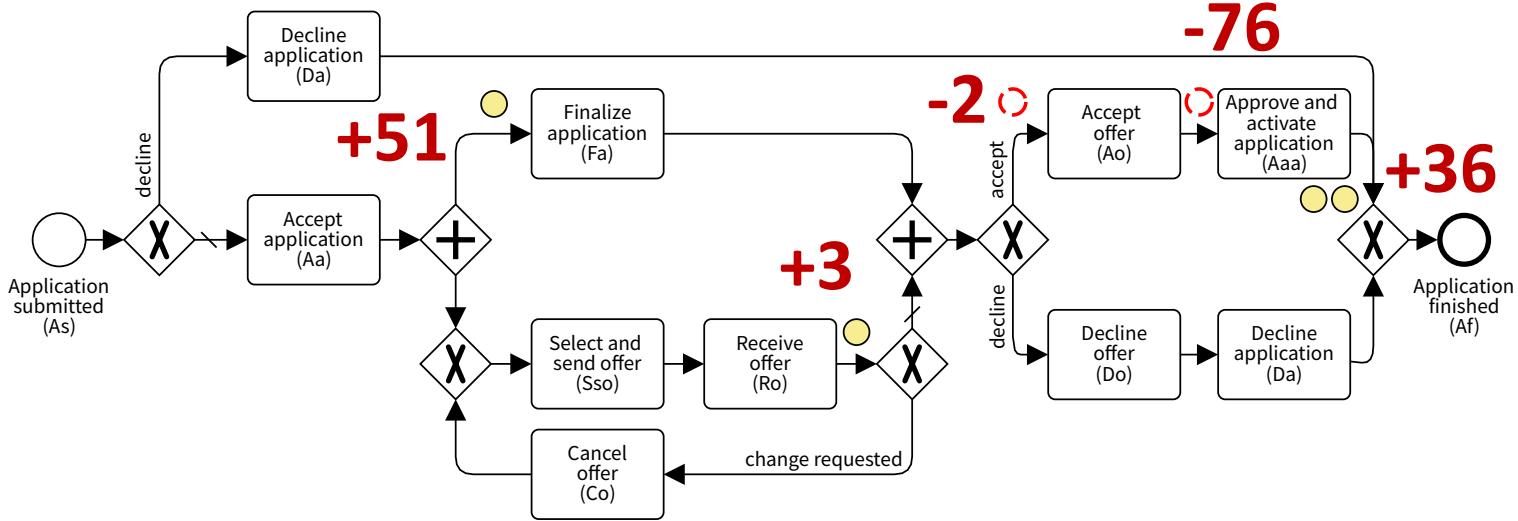
Observed Trace: < As, Aa, Sso, Ro, Ao, Aaa, Aaa >

Local Feedback on Non-Conformance



Deviations \approx Tokens Missing or Tokens Remaining after replaying a trace

Aggregated Feedback on Non-Conformance



- Aggregated information about missing and remaining tokens
- Identification of “hotspots” of non-conformance
- Highlights major issues when replaying log and separates some from noise

Fitness Measure

Determine ratios based on missing and remaining tokens:

$$\text{fitness}(\sigma, M) = \frac{1}{2} \left(1 - \frac{\text{missing}(\sigma, M)}{\text{consumed}(\sigma, M)} \right) + \frac{1}{2} \left(1 - \frac{\text{remaining}(\sigma, M)}{\text{produced}(\sigma, M)} \right)$$

$$\text{fitness}(L, M) = \frac{1}{2} \left(1 - \frac{\sum_{\sigma \in L} \text{missing}(\sigma, M)}{\sum_{\sigma \in L} \text{consumed}(\sigma, M)} \right) + \frac{1}{2} \left(1 - \frac{\sum_{\sigma \in L} \text{remaining}(\sigma, M)}{\sum_{\sigma \in L} \text{produced}(\sigma, M)} \right)$$

Overall conformance assessment based on aggregated fitness measure

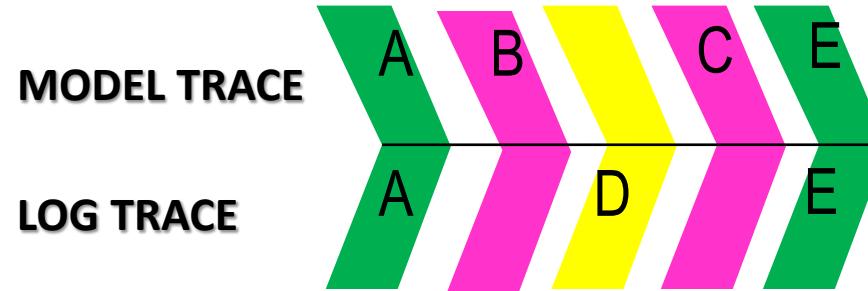
Condenses (non-)conformance into a single value

Yet, hard to interpret in terms of the absolute value

Computing Conformance Checking Artefacts

Alignments

Conformance Artefact: Alignments

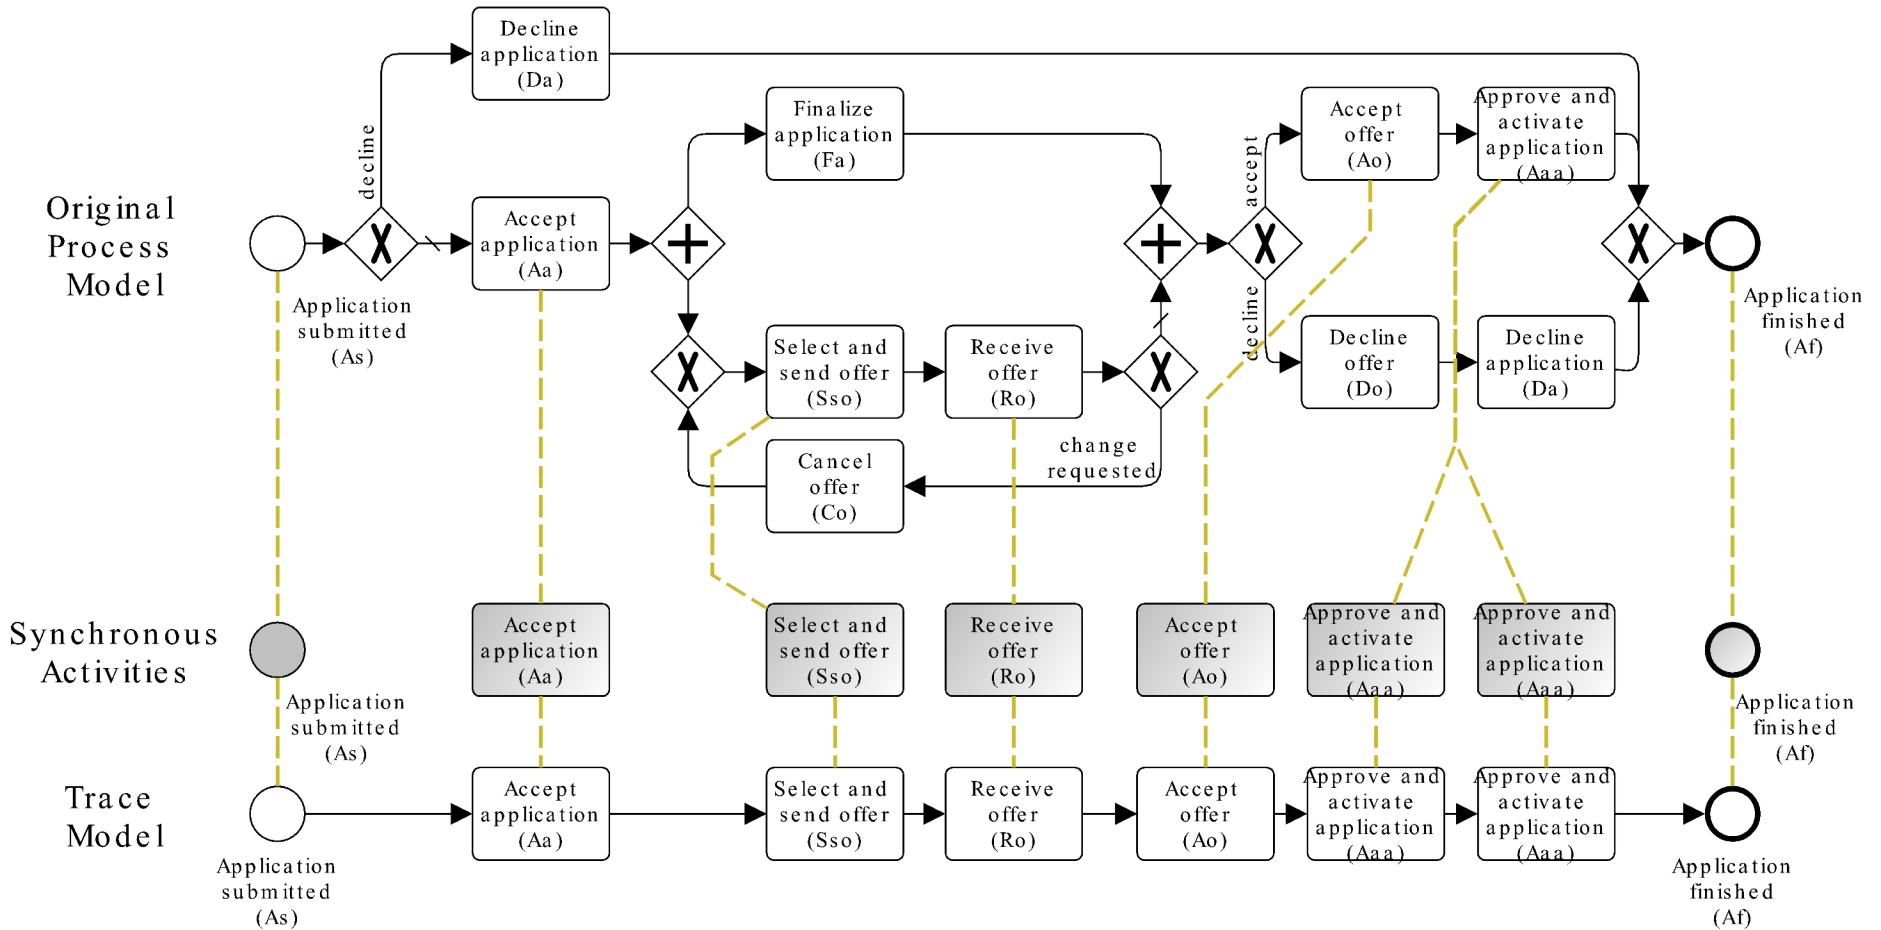


- **SYNCHRONOUS MOVE: LOG AND MODEL AGREE**
- **MODEL MOVE: TASKS NOT RECORDED THAT SHOULD BE EXECUTED**
- **LOG MOVE: EVENTS RECORDED THAT CANNOT BE EXECUTED**

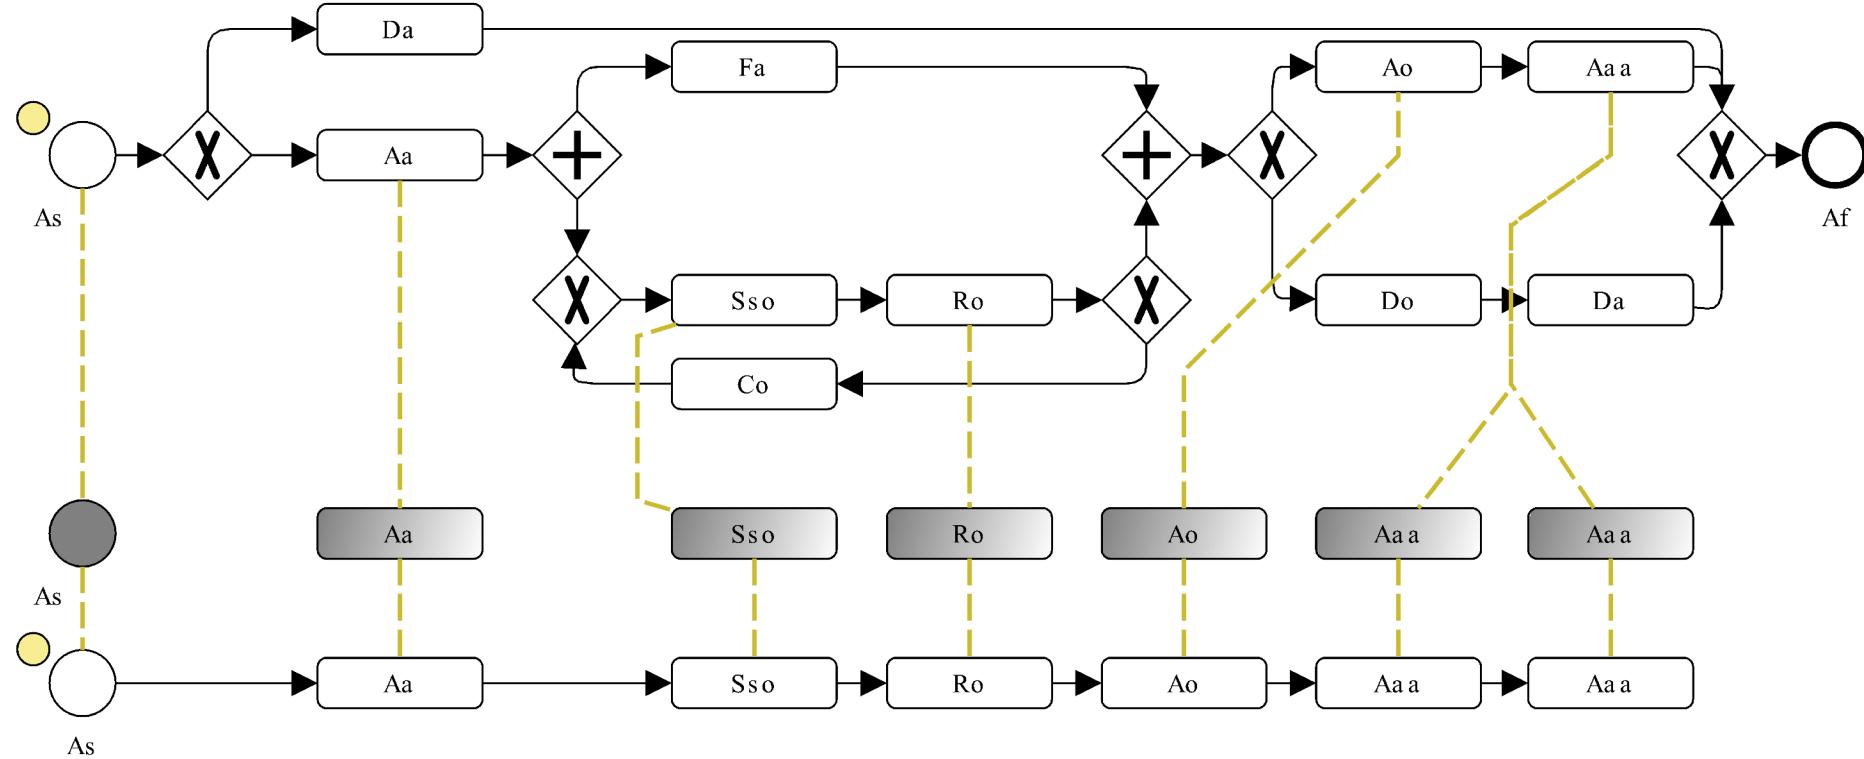
Computing Alignments through A*

- Reference technique for computing alignments
- Can guarantee **optimality** for an arbitrary cost function on deviations
- Can be adjusted to find not just one but **all** optimal alignments
- Implemented in ProM, with many optimizations.
- Based on the notion of *Synchronous Product Net*

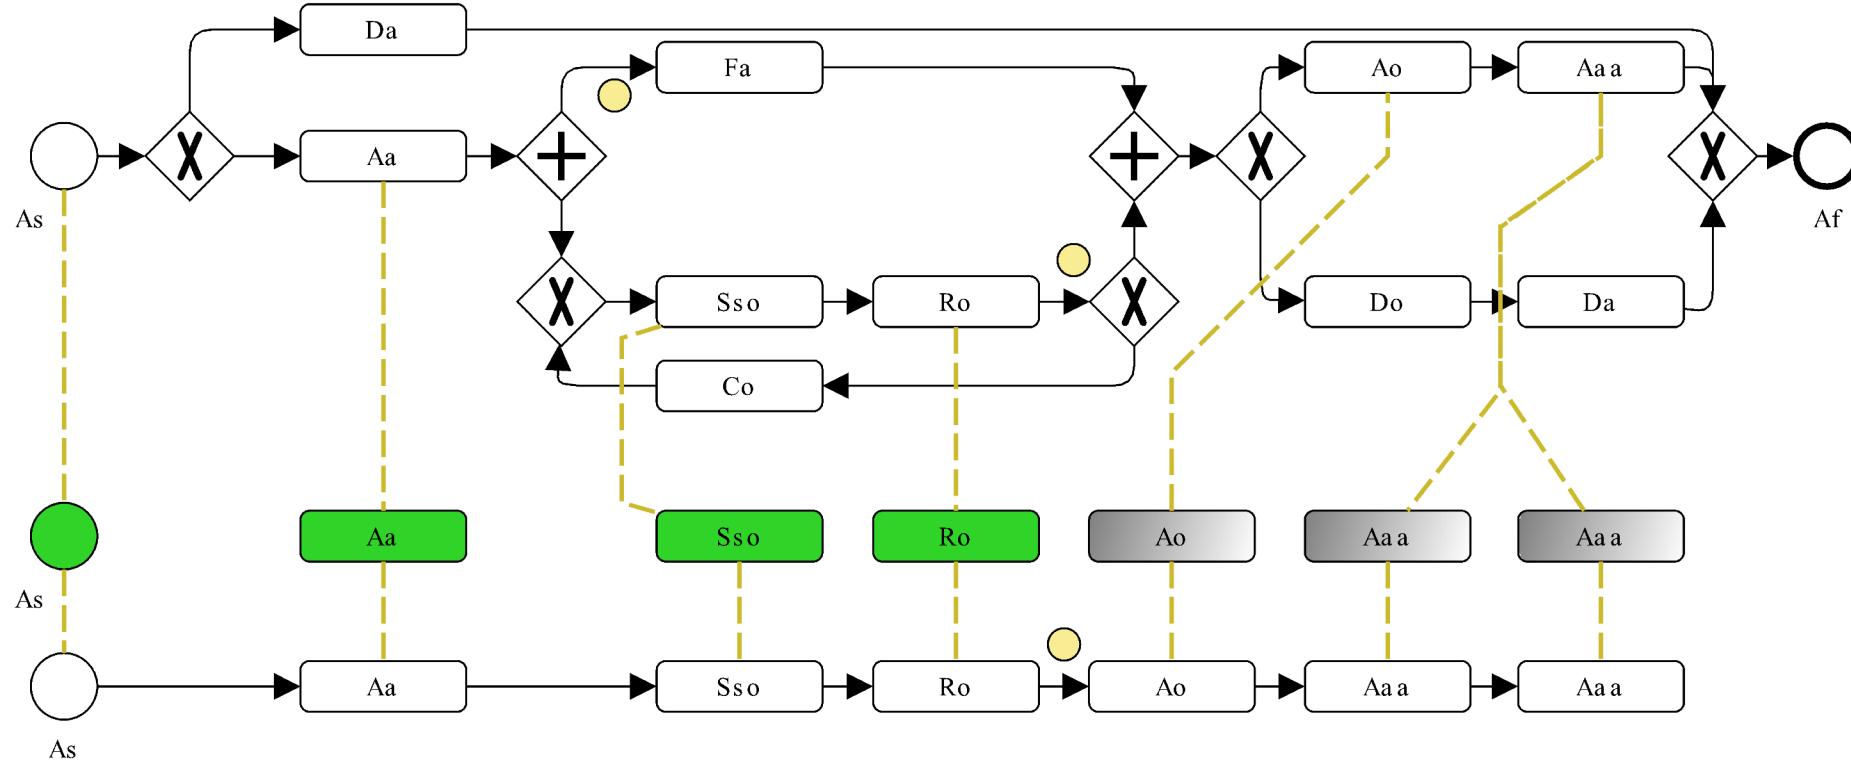
Synchronous Product Net (SPN)



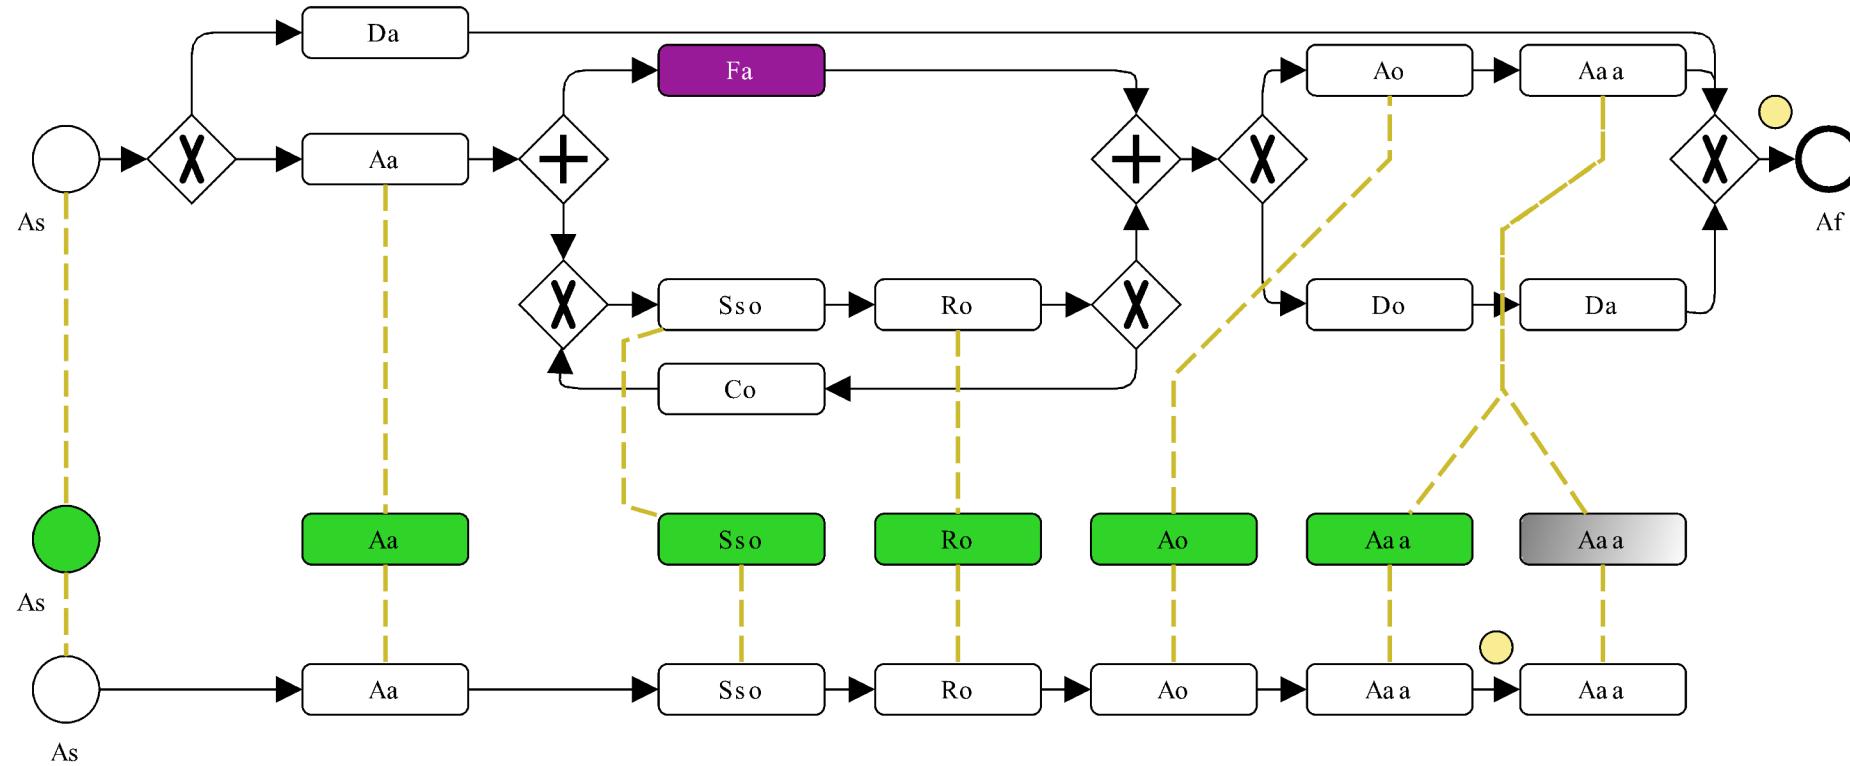
A run of the SPN



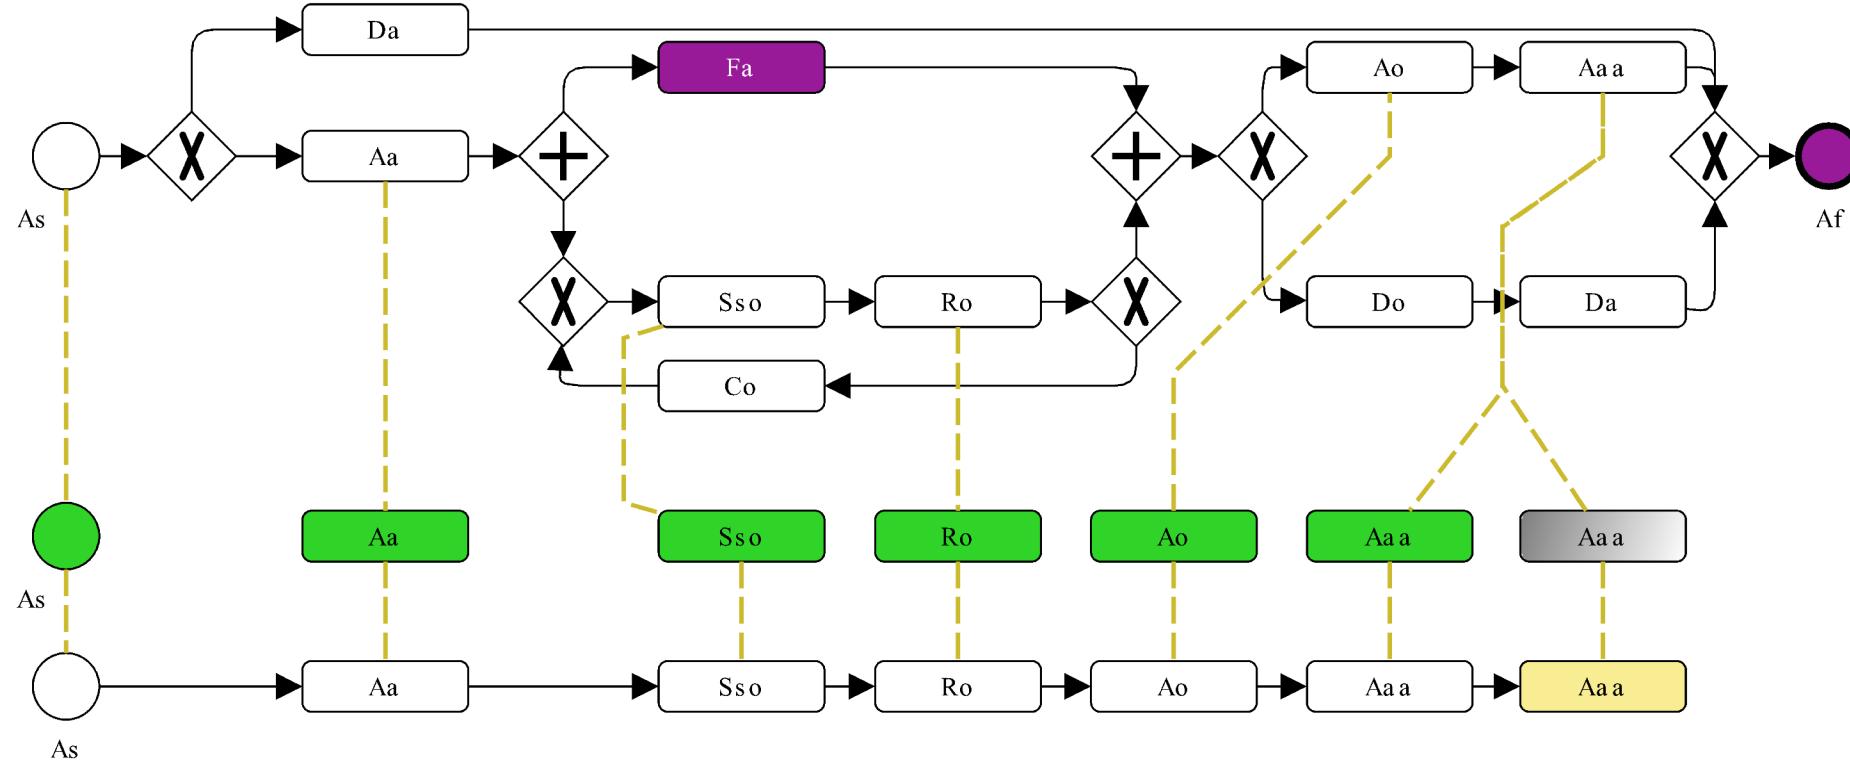
A run of the SPN



A run of the SPN



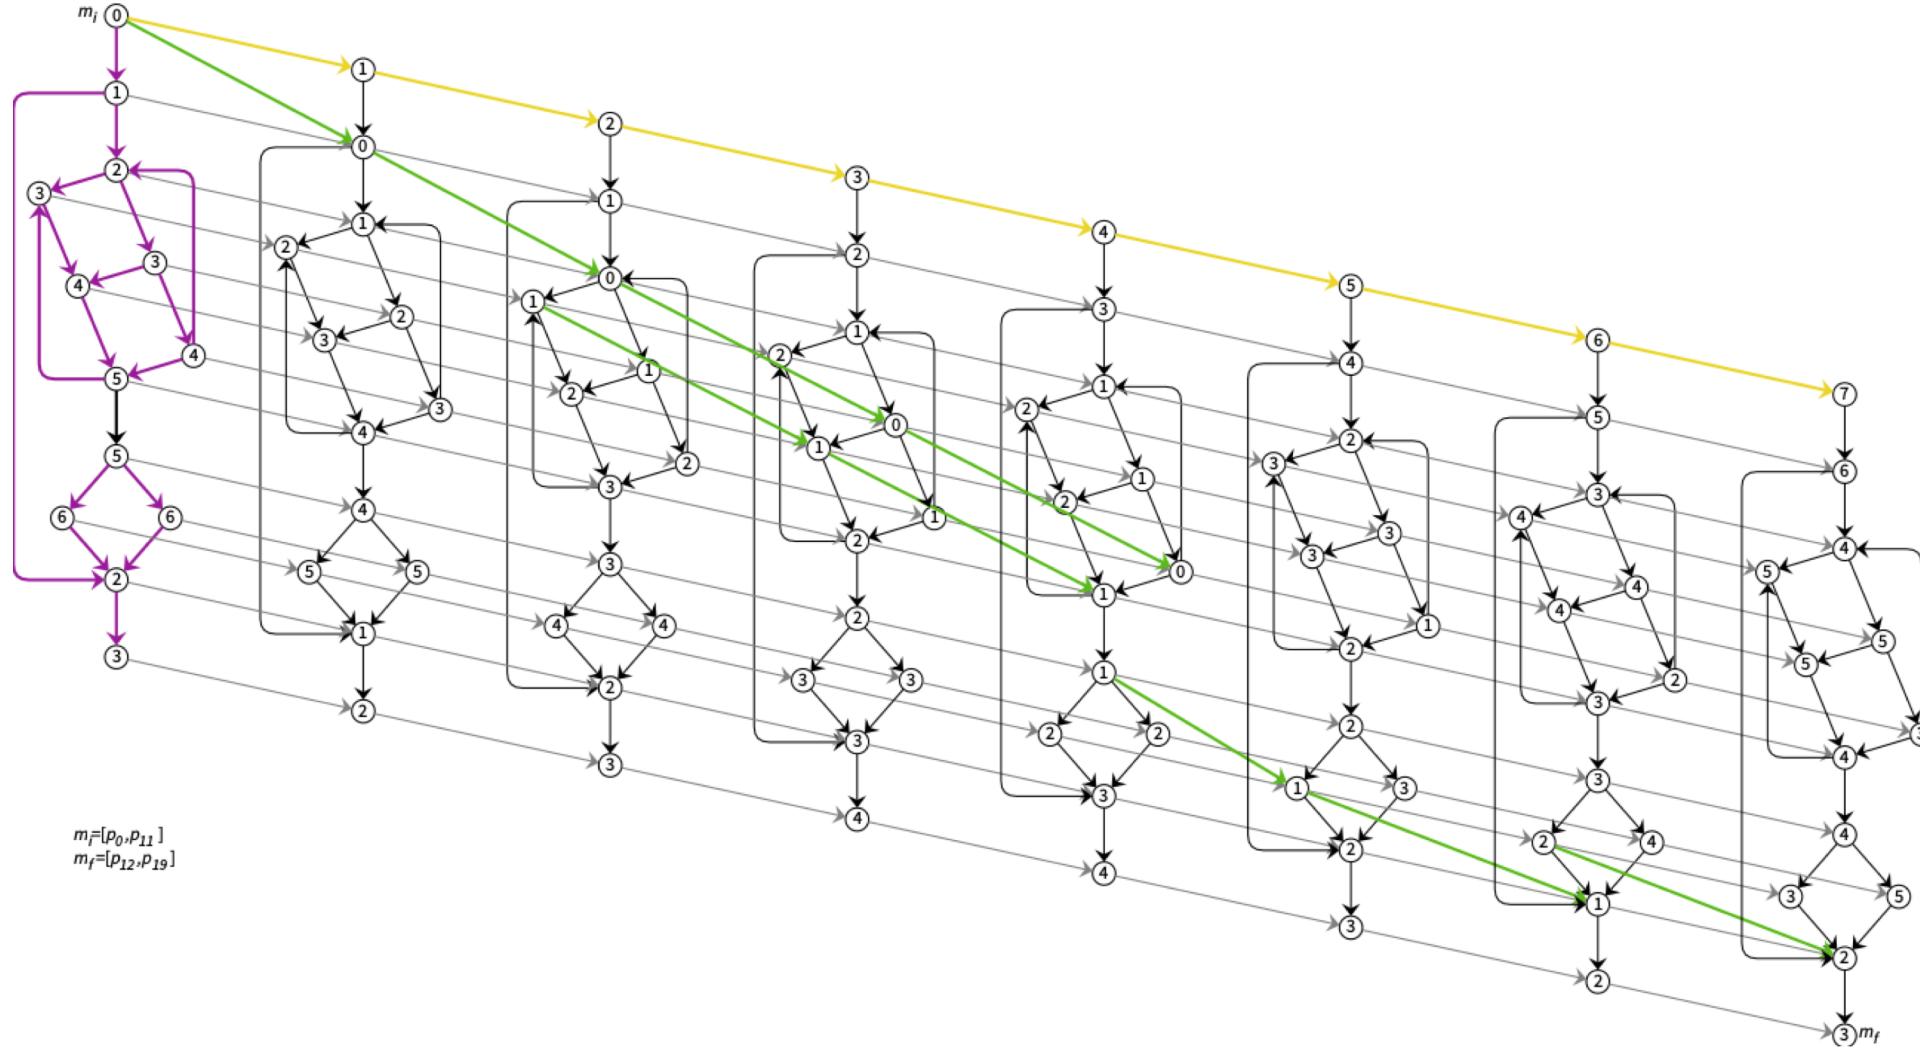
A run of the SPN



Algorithmics Behind the A*

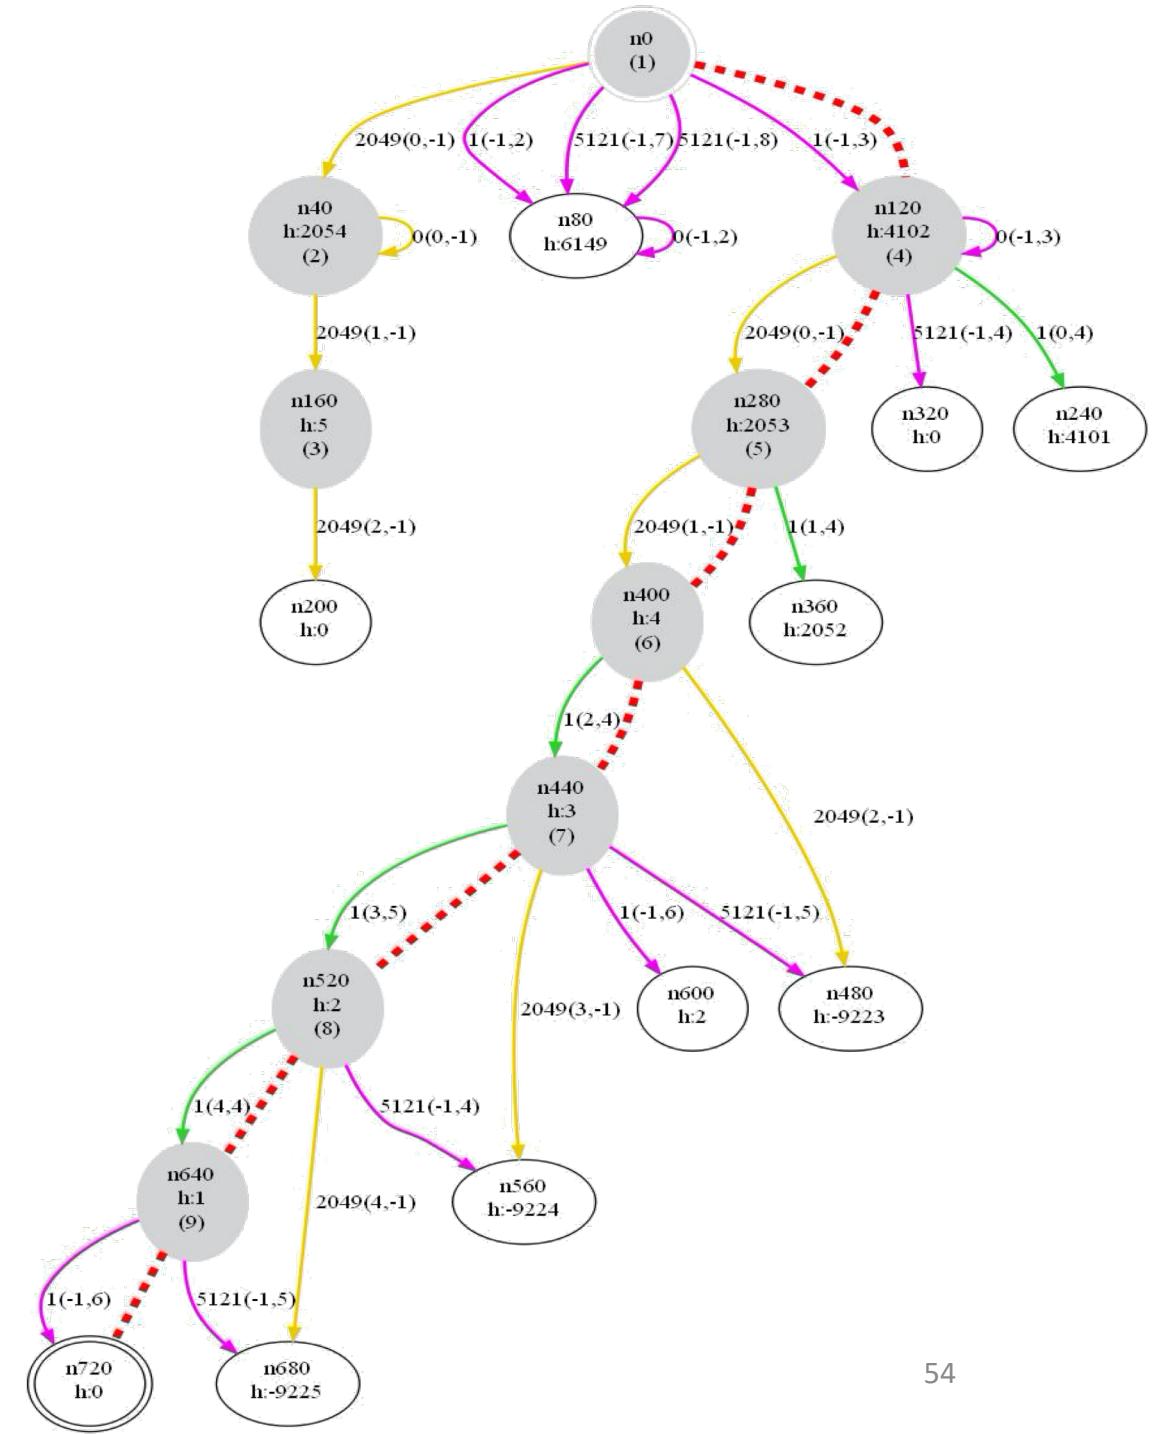
- In general, the SPN has an infinite set of runs !
- Instead: Traverse the SPN **state space** in a efficient way so that a cost-minimal path from the initial to the final state (representing the optimal alignment) is found.
- Using an estimation of the distance to the final state, prune parts of the state space.

SPN State Space



Computing Alignments

- The search space is the statespace of the synchronous product model
- Each node is a combination of a state in the model and the remaining events in the trace
- Each arc is a **move on model**, **move on log** or a **synchronous move**
- A heuristic function estimates the remaining distance to the final node (i.e. model in a final state and all events executed)
- We find a shortest path!



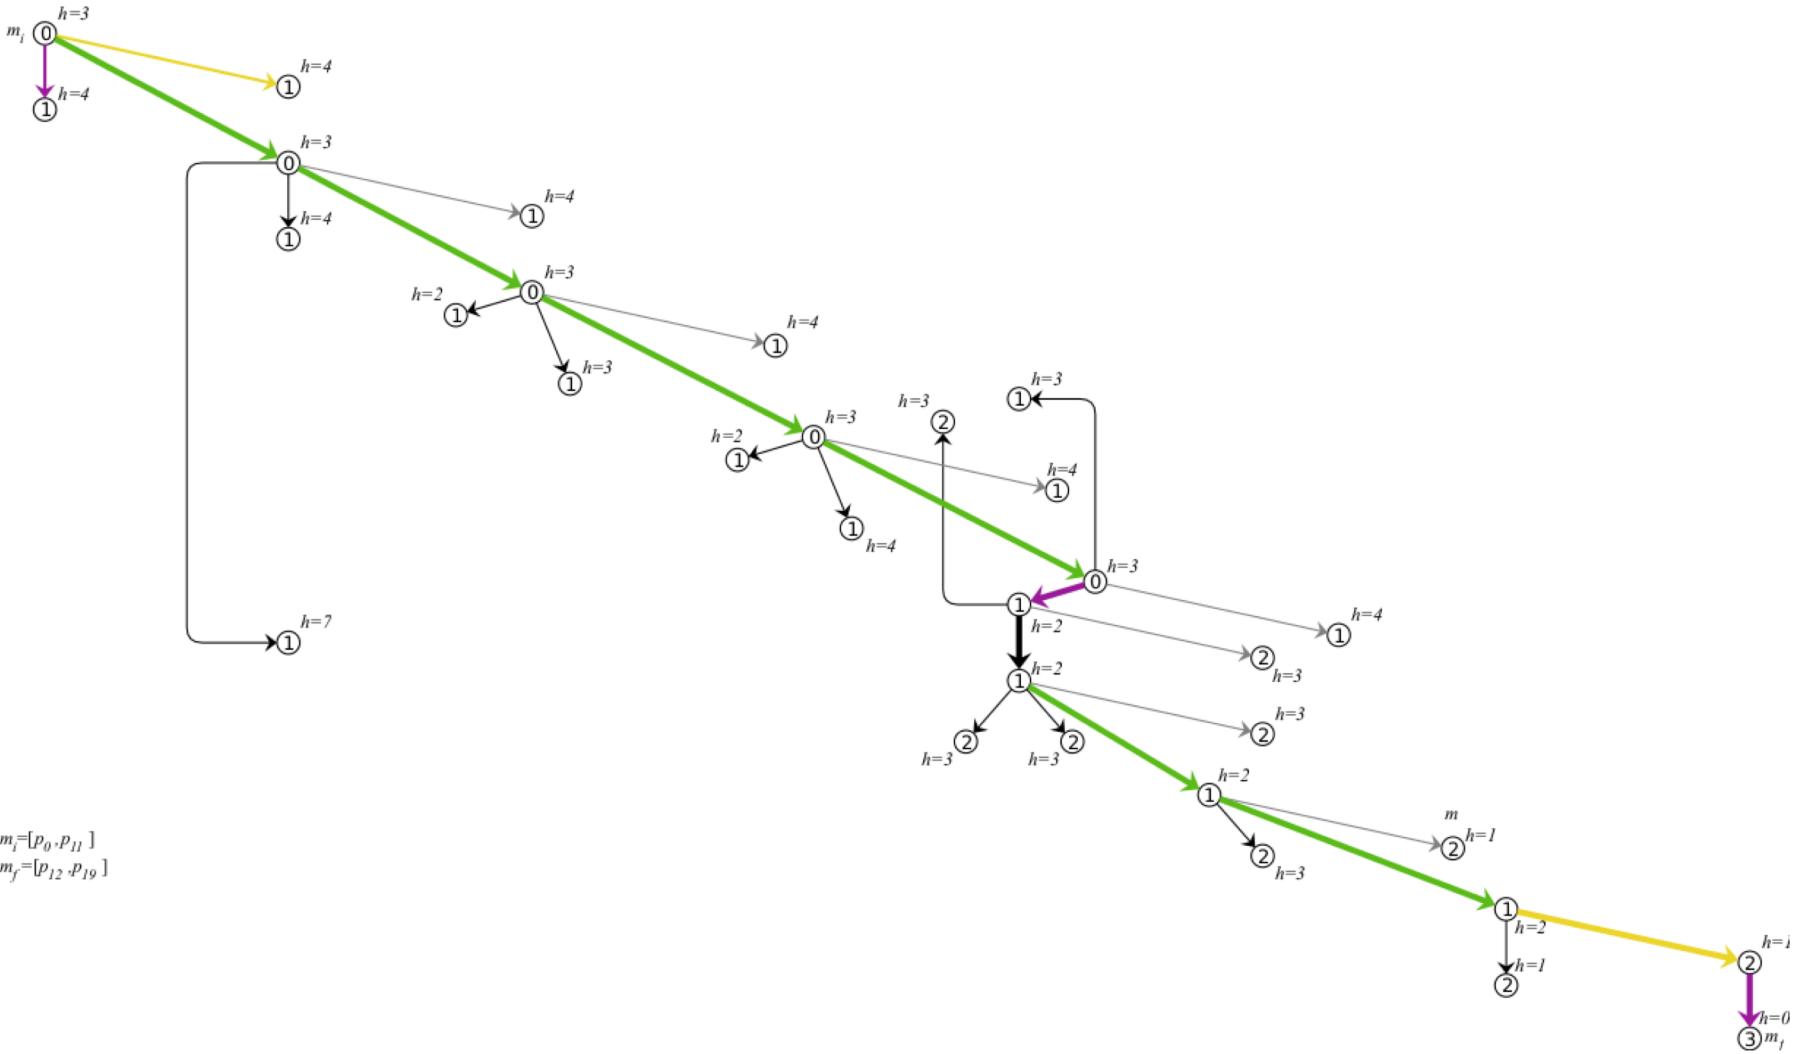
Computing Alignments (A* with fast lowerbound)

```
Initialize HashBackedPriorityQueue q
While (peek(q) is not target t)
    VisitedNode n = head(q)
    If the estimate for node(n) is not exact
        Compute the exact estimate for the remaining distance to t and
        add n to the priority queue
    continue
    add n to the considered nodes
    For each edge in the graph from node(n) to m
        If m was considered before, continue
        If m is in the queue with lower cost, continue
        If m is in the queue with higher cost, update and reposition it
        If m is new,
            compute a fast lowerbound for the remaining distance to t
            v = new VisitedNode(m)
            set n a predecessor for v
            add v to the priority queue
    Return head(q)
```

The diagram illustrates the time complexities for various steps in the A* algorithm:

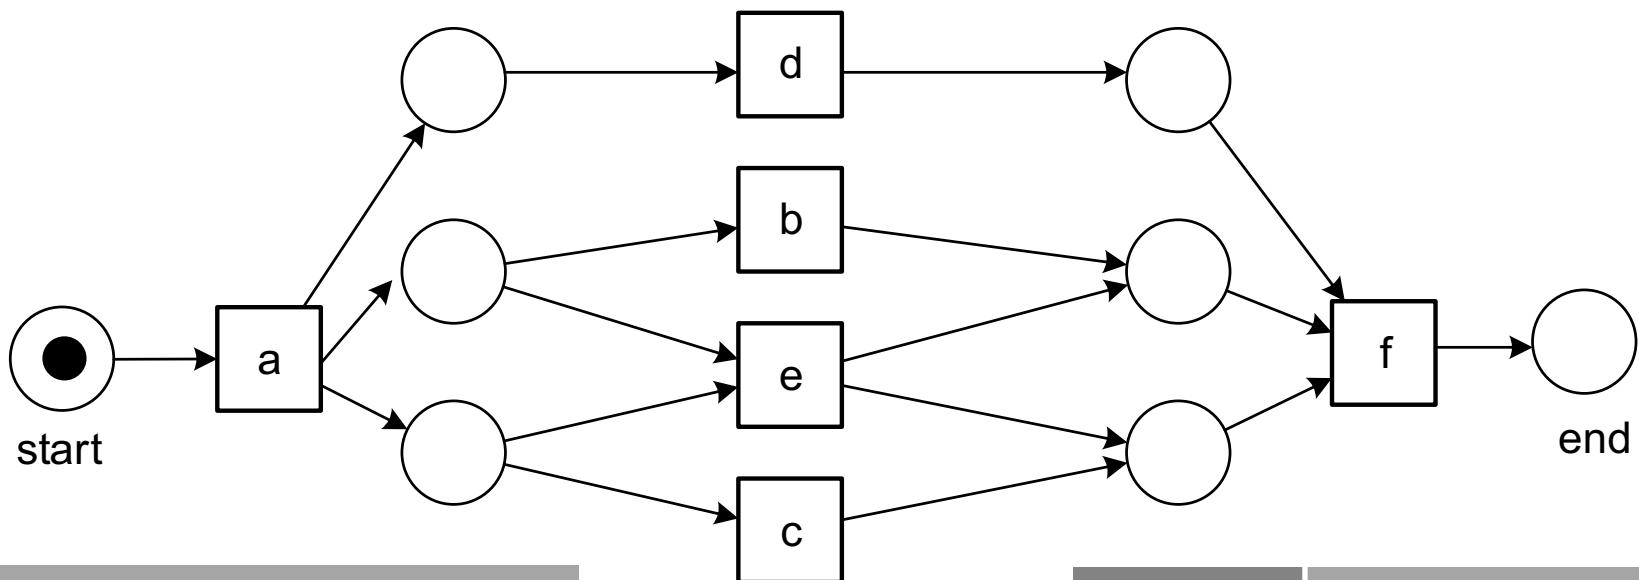
- Initial step: $O(\log(\text{size } q) + \alpha)$
- While loop condition: $?$
- VisitedNode extraction: $O(\log(\text{size } q) + \alpha)$
- Exact estimate computation and enqueue: $O(\log(\text{size } q) + \alpha)$
- Considered nodes update: $O(\log(\text{size } q) + \alpha)$
- Edge processing loop:
 - Edge iteration: $O(\log(\text{size } q) + \alpha)$
 - Conditionals (lower cost, higher cost, new): $O(1)$
- Final step: $O(\log(\text{size } q) + \alpha)$

SPN State Space after A* Prunning



Alignments

- Find the optimal alignment(s).



trace	a	e	c	d	>>
model	a	e	>>	d	f

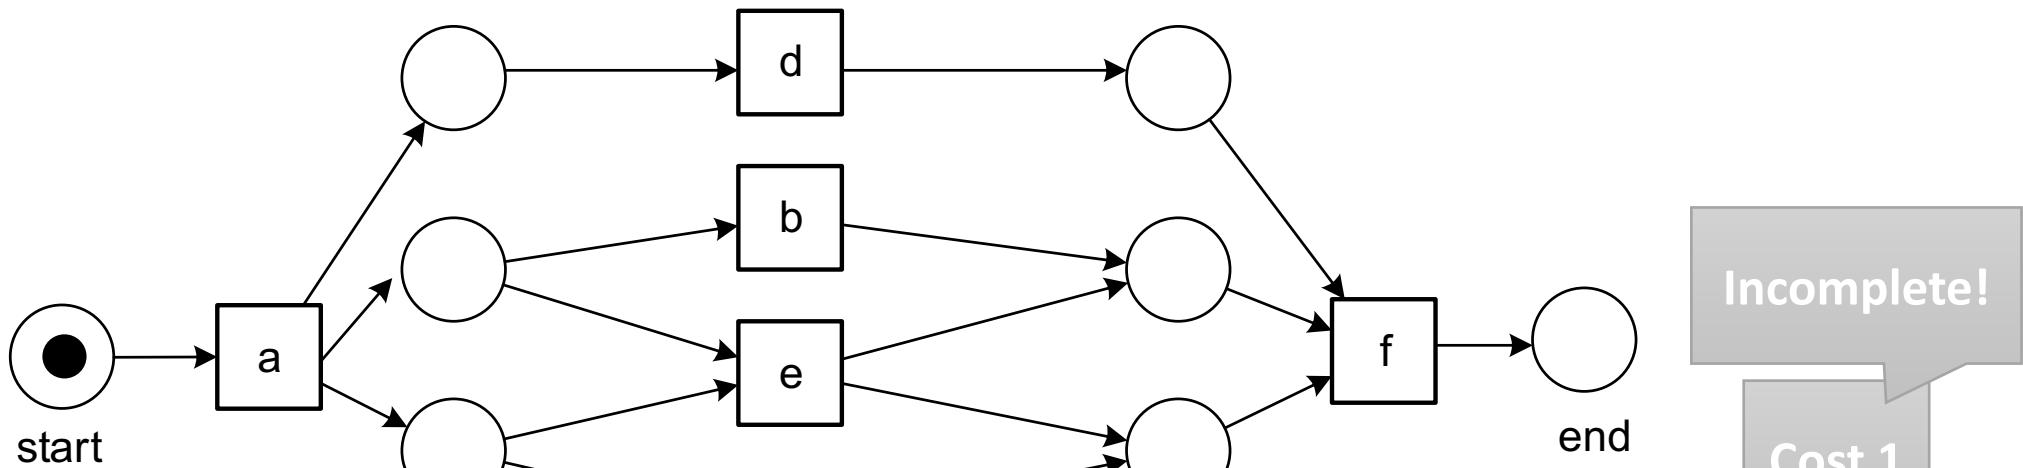
trace	a	e	c	d	
model	a		>>	c	d

trace	a	e	c	d	>>	>>
model	a	>>	c	d	b	f

trace	a	e	c	d	>>	>>
model	a	e	>>	>>	d	f

Alignments - solution

- Find the optimal alignment(s).



Cost 6

trace	a	e	c	d	>>
model	a	e	>>	d	f

trace	a	e	c	d	>>	>>
model	a	>>	c	d	b	f

Incomplete!

Cost 1

trace	a	e	c	d	>>
model	a		>>	c	d

trace	a	e	c	d	>>	>>
model	a	e	>>	>>	d	f

Cost
12

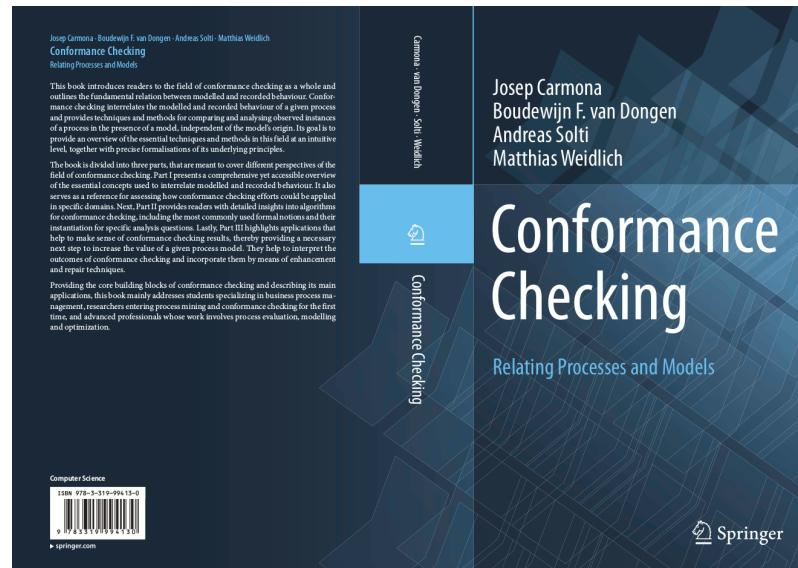
"Conformance Checking: Relating Processes and Models"

Josep Carmona

Boudewijn van Dongen

Andreas Solti

Matthias Weidlich



<https://www.springer.com/gb/book/9783319994130>